

# Software-implemented hardening against soft errors

Massimo VIOLANTE

Politecnico di Torino - Italy

# Goal

- To present an overview of soft error mitigation techniques based “mainly” on software approaches
- In this talk we will look at the so-called **Software-Implemented Hardware Fault Tolerance (SIHFT) Techniques**



# Outline

- Introduction: why SIHFT?
- Assumption
- A system-level view of soft errors
- SIHFT target
- SIHFT techniques
- Conclusions



# Introduction: why SIHFT?

- COTS components are growing interest due to performance reason, but
- COTS components are not intended for tolerating soft errors, and must be protected
- SIHFT provides soft error mitigation using time and information redundancies in software to avoid hardware redundancy
  - More hardware = more expenses
  - More software ≠ more expenses, iff robust software can be obtained automatically



# Assumptions

- **Hardware:**

- Microprocessor-based systems
- COTS components (not rad hard)
- Single Event Upset (SEU)

- **Software:**

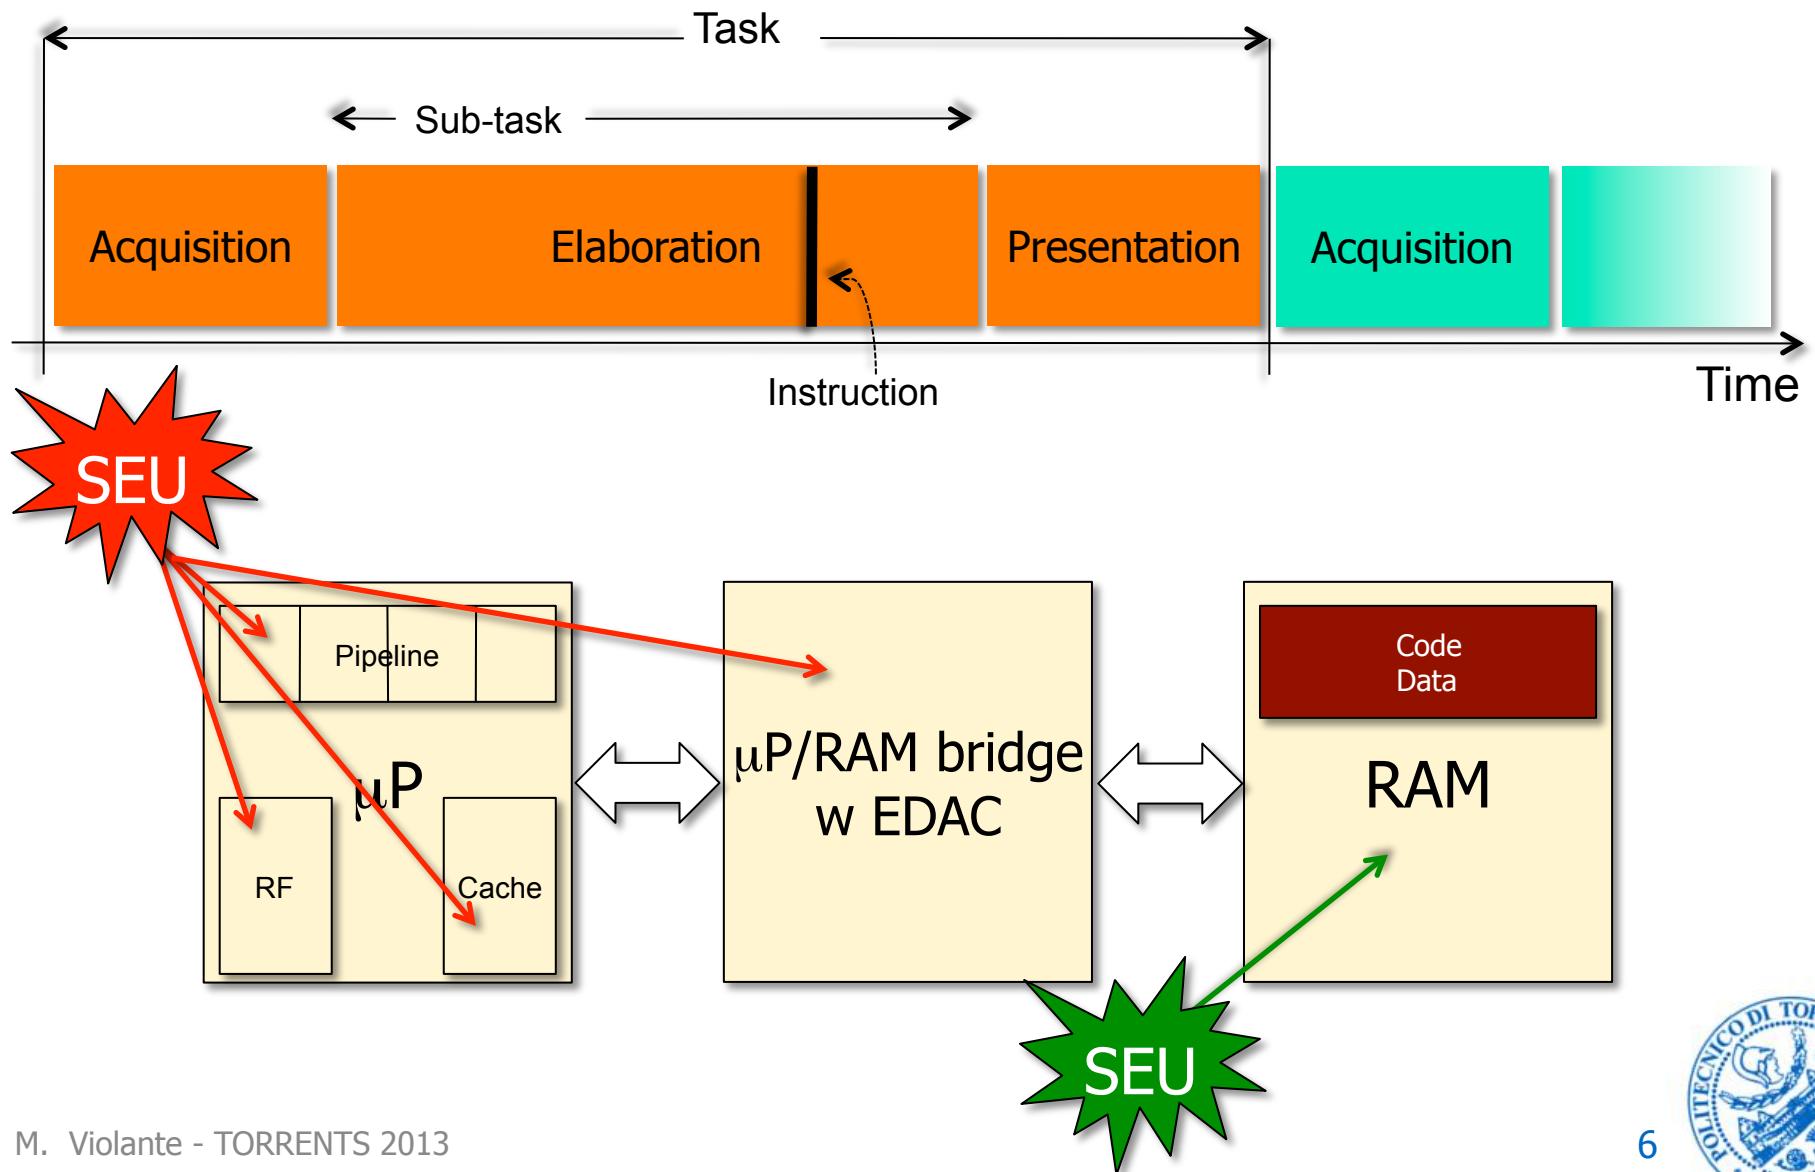
- It is correct (no bugs)
- It is fully available as source code

- **Target of the system:**

- Payload processing (one task, no operating system)



# Assumptions (cont.)



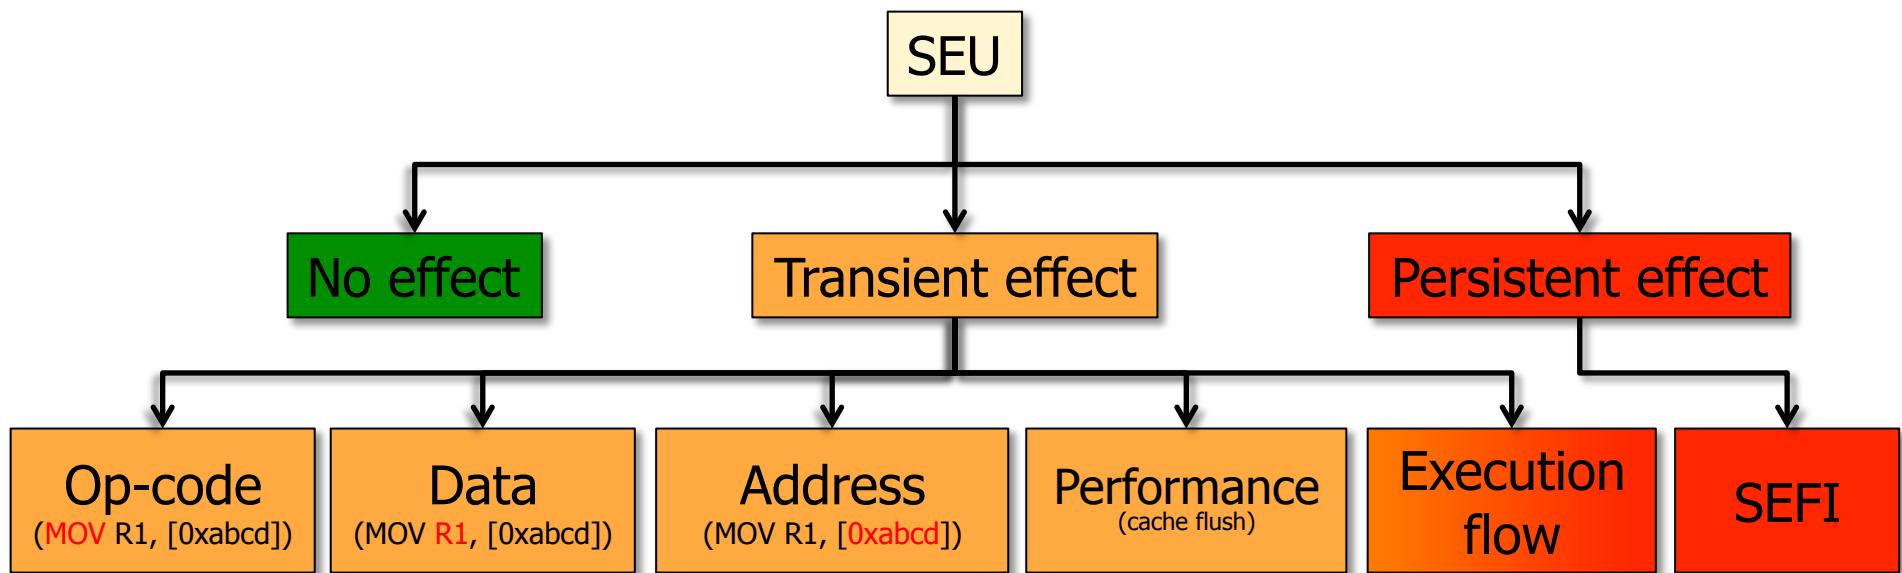
# Assumptions (cont.)

- Active redundancy is enforced:
  - The fault is detected first, then
  - The fault is corrected
  - No fault masking (like TMR)
- Correction can be:
  - Forward correction
    - Correct outputs generated from faulty outputs using redundant data
  - Backward correction
    - The computation is repeated



# A system-level view of soft errors

- SEU effects can be modeled according to the modifications they introduce to a running program
  - Model does not depend on the SEU location

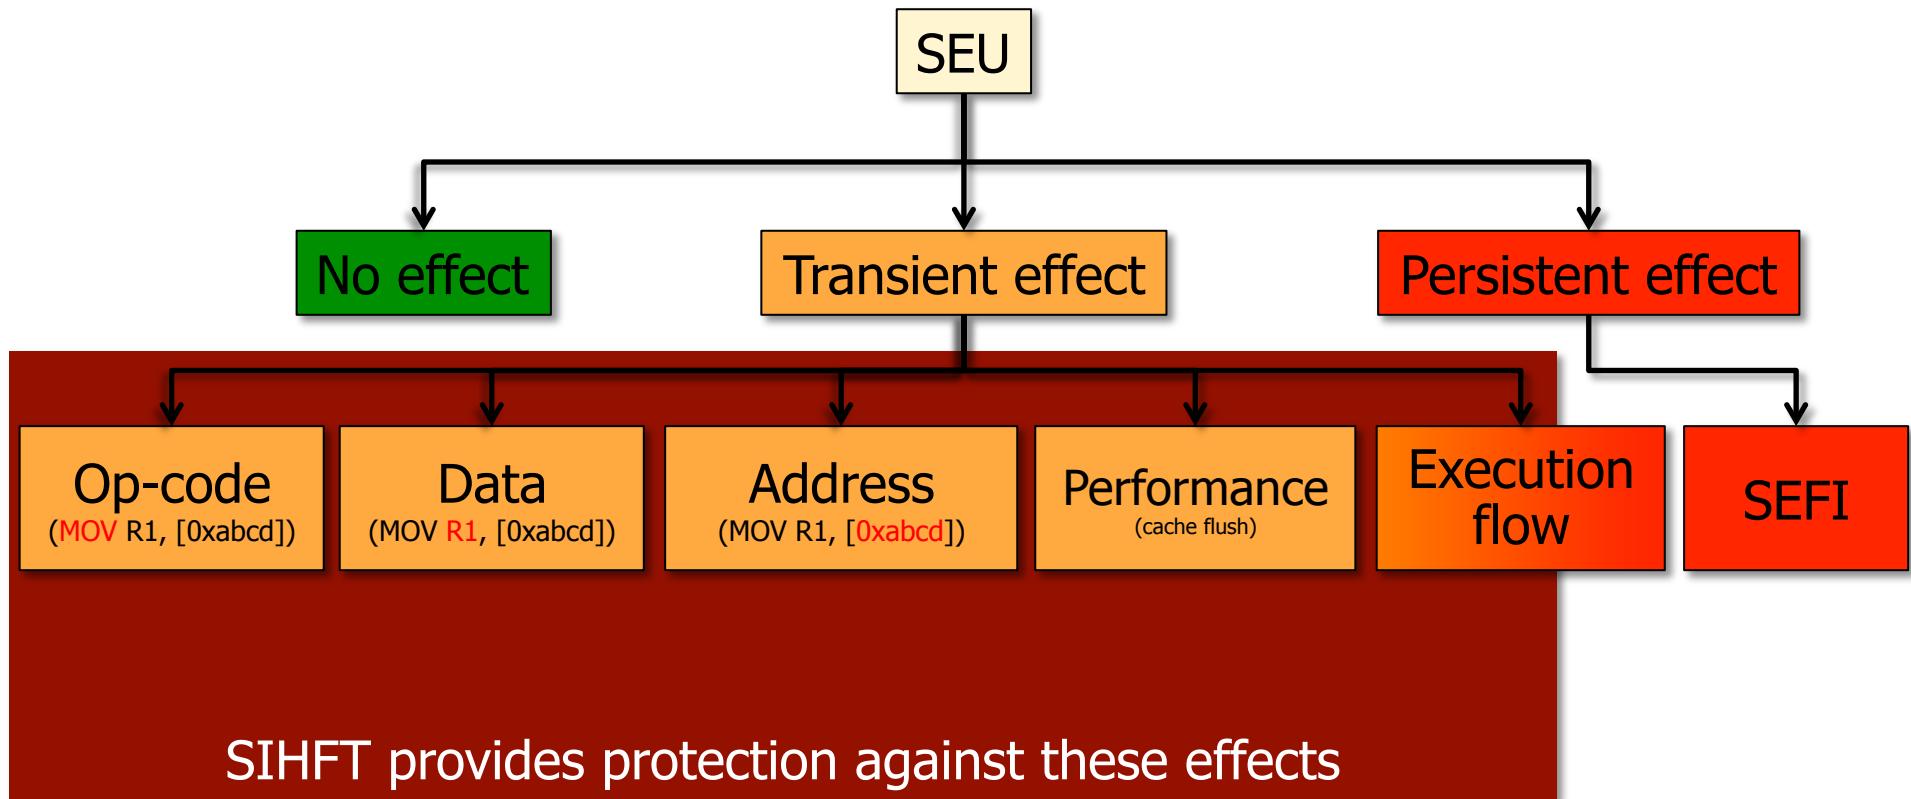


# SIHFT target

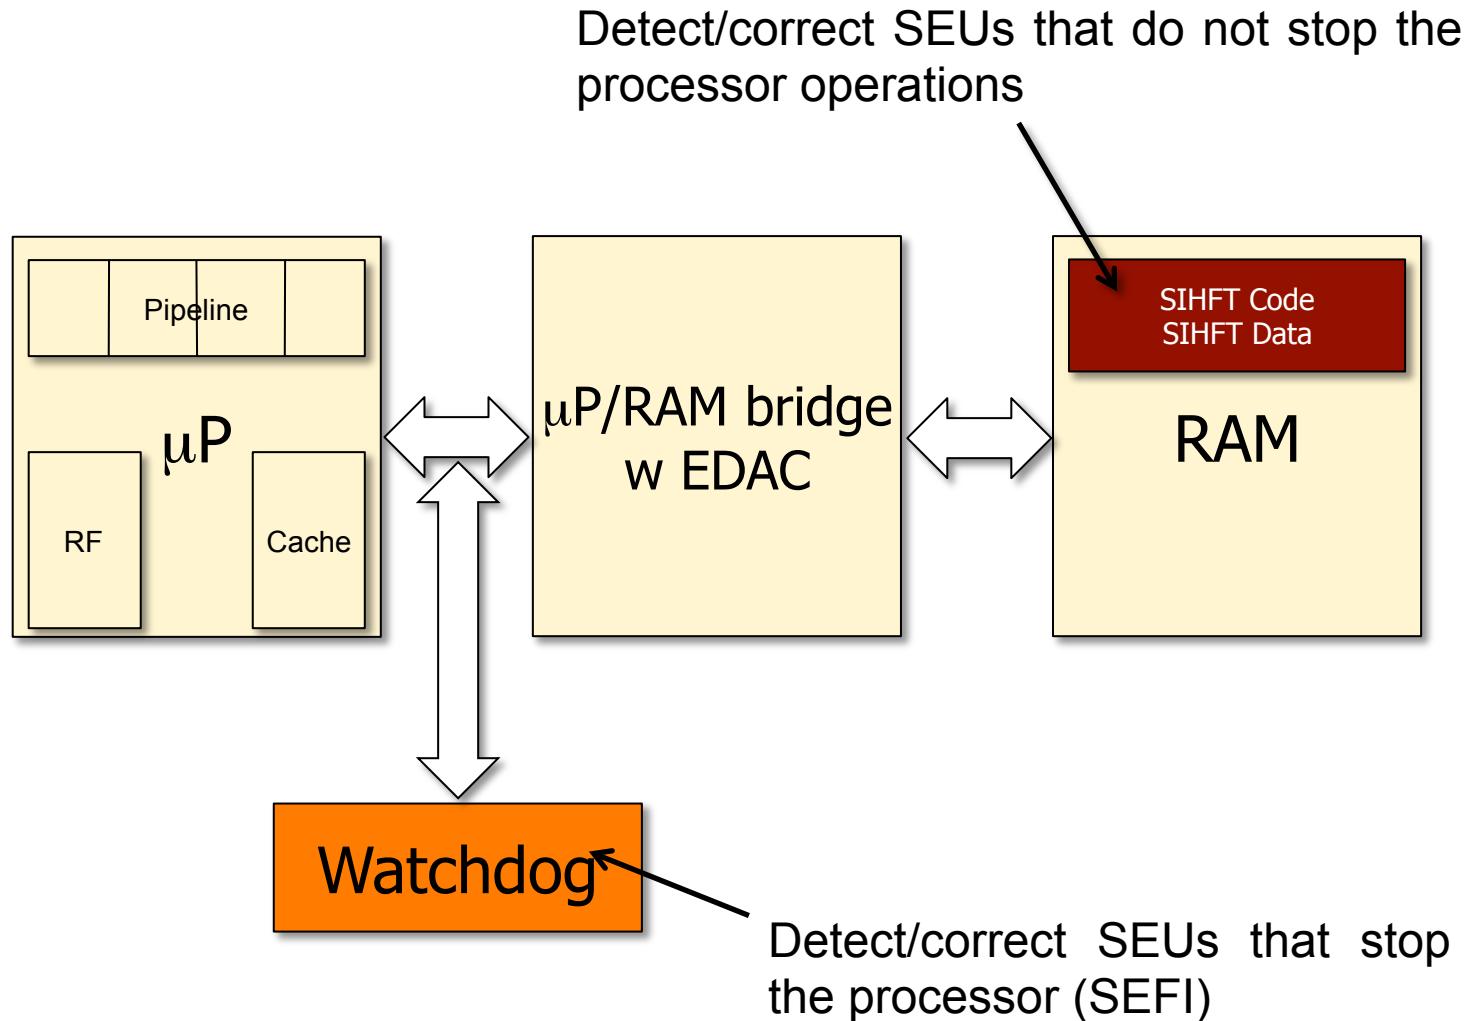
- SIHFT techniques target SEU effects that does not “hang” the processor (SEFI)
  - Being based on software, SIHFT mandates a processor that is up and running
- To cope with SEFI, complementary techniques are needed
  - Low-cost hardware working in parallel with main processor (e.g., watchdog)



# SIHFT target (cont.)



# System architecture



# SIHFT techniques

- Instruction-level time redundancy
  - Data-oriented techniques
  - Control-oriented techniques
- Task-level time redundancy



# SIHFT techniques

- Instruction-level time redundancy
  - Data-oriented techniques
  - Control-oriented techniques
- Task-level time redundancy



# Instruction-level time redundancy

- Based on adding instructions (at C level, assembly level, or intermediate level) to:
  - Replicate data
  - Replicate computations
  - Perform consistency checks
- Two classes of techniques:
  - Data-oriented for tackling everything but errors in the execution flow
  - Control-oriented for tackling execution-flow errors
- Exploit backward error correction
- SEFI are not detected



# Data-oriented techniques

- The source code of application software is modified by:
  - Replicating each stored data
  - Replicating each operation
  - Checking the consistency of data
- Can be applied at different levels:
  - High-level (e.g., C) programs
  - Assembly programs
  - On compiler-generated intermediate code



# EDDI

- Proposed in 2002 by Oh, Shirvani and McCluskey
- Based on duplicating instructions and data at the assembly level
- Reduces execution time slow down by careful instruction scheduling in superscalar architectures



# Example

- Original code

```
ADD R3, R1, R2
```

- Modified code

```
ADD R3, R1, R2  
ADD R23, R21, R22  
BNE R3, R23, error
```



# EDDI results

- Detection capability 100%
- Execution time overhead ranges from 72% to 111% for a two-way processor
- Memory overhead is about 100%



# The PoliTo approach

- Proposed by Politecnico di Torino group starting from 1999
- Main characteristics:
  - Works on high-level code (e.g., C language)
  - Limited assumptions on addressed faults
  - High independence on hardware
  - High fault coverage
  - Its application can be automated
  - Originally intended for fault detection, only



# Basic idea

- A set of rules have been defined
- Rules transform a high-level code into a hardened one
- Rules application can be automated



# Rules

- Duplicate every variable
- Execute every operation on the two replicas
- Check for consistency after each read access



# Example

- Original code

```
int a, b;  
...  
b = a+5;
```

- Modified code

```
int a1, b1, a2, b2;  
...  
b1 = a1+5;          b2 = a2+5;  
assert( a1==a2 );
```



# Transformation tool

- Transformation rules can be applied automatically
- A prototypical tool automating the application of the rules has been developed
  - Reads a C code
  - Produces a hardened C code



# Overhead

- The application of the hardening rules on the whole code on unpipelined or RISC processor systems
  - Increases the size of the code and data areas
    - Factors **from 3 to 4** have been observed
  - Reduces the program execution
    - Factors of **about 3** have been observed
- Main issue
  - The need for disabling compiler optimizations, or
  - Adoption of ad-hoc compiler



# Experimental results

- Several fault injection campaigns have been performed to evaluate the fault detection capabilities of the method
  - On a transputer-based system
  - On an 8051-based system
  - On a LEON processor



# Experimental set up

- Some sample benchmark programs have been selected
- Hardened versions have been developed
- Fault injection and radiation experiments have been performed



# Fault Classification

- Effect less
- No Answer
- Latent
- Wrong Answer
  - Undetected incorrect output
- Detected

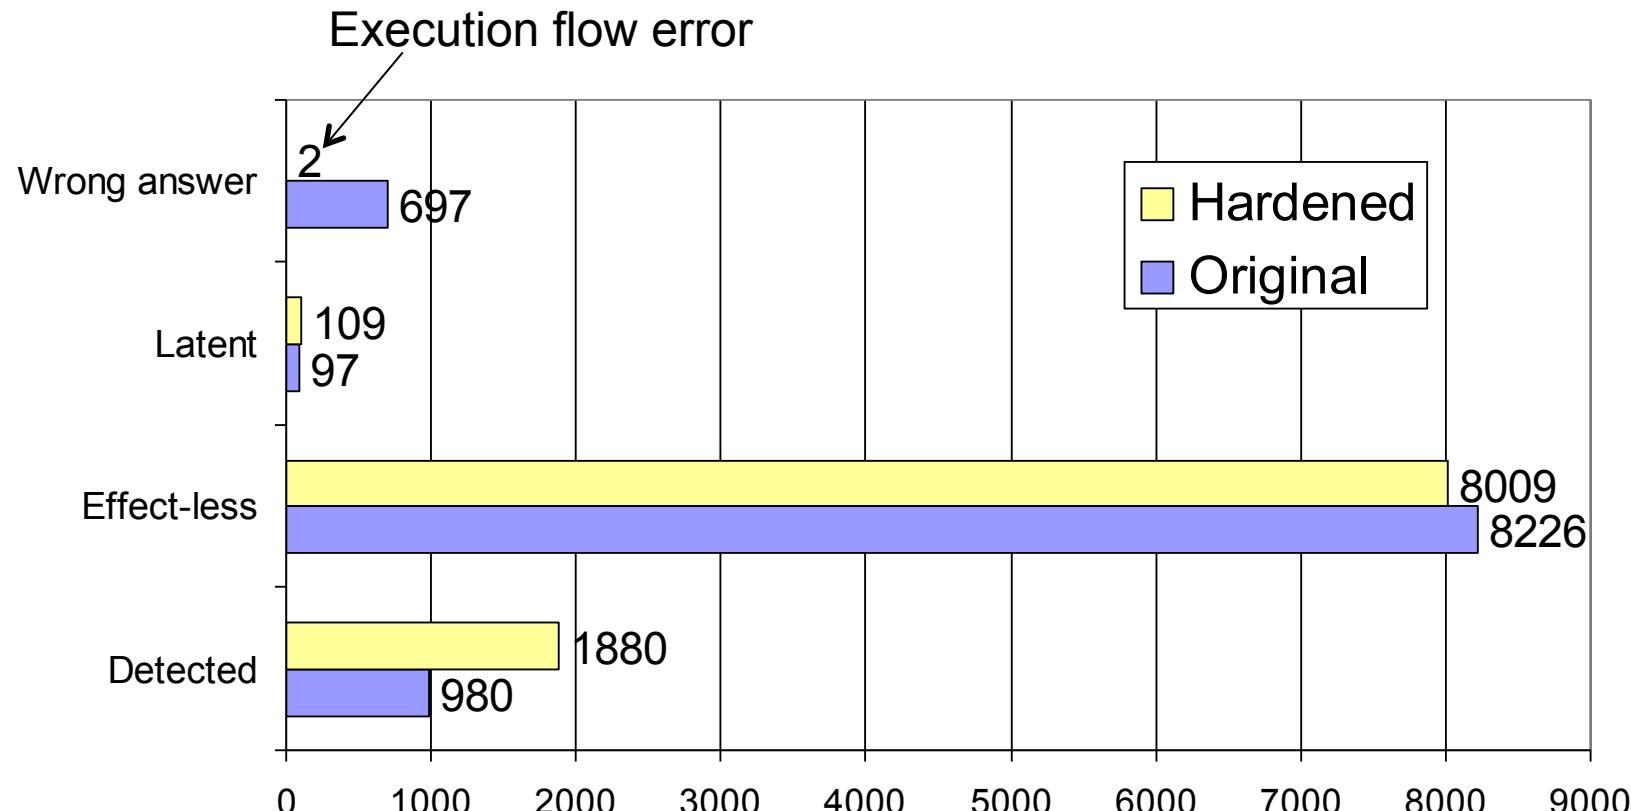


# LEON system: set up

- Injected faults
  - SEUs in the processor memory elements (including register file and pipeline registers)
- Emulation-based fault injection has been exploited

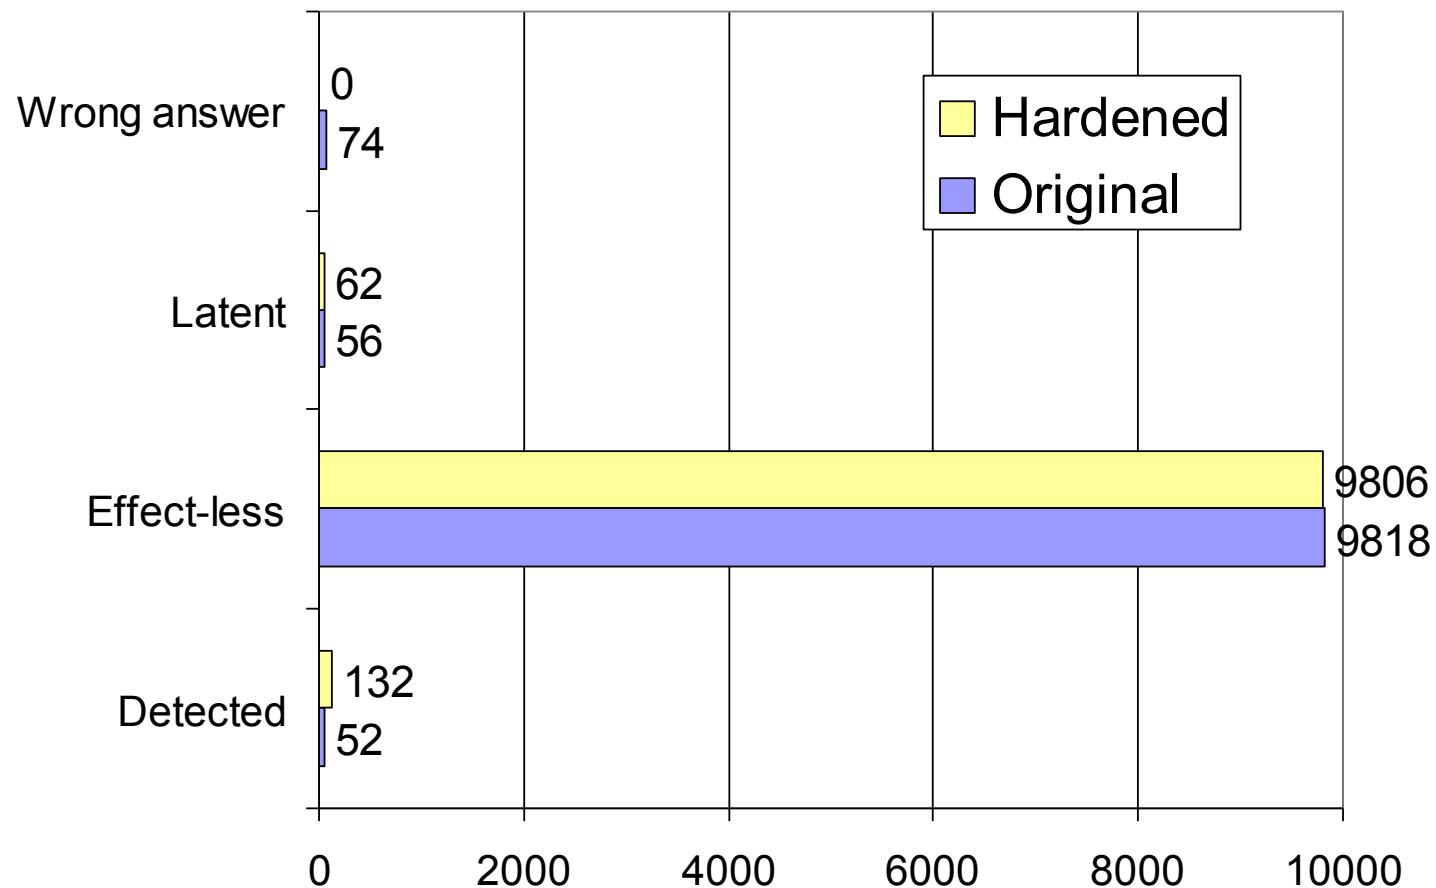


# SEUs in pipeline registers



10,000 SEUs injected

# SEUs in register file



10,000 SEUs injected

# Observations (I)

- The method detects faults in:
  - Cache
  - User registers
  - Hidden registers (e.g., in the control unit, or in the pipeline)
  - Combinational logic



# Observations (II)

- The method is able to detect:
  - Any kind of fault creating a mismatch between the two replicas of a variable
  - Most of the faults affecting executed instructions
- The method is **NOT** able to detect:
  - Faults affecting the execution flow
  - Persistent faults



# Observations (III)

- The method can be applied **flexibly** :
  - A subset of rules may be applied (e.g., only duplication rules)
  - A subset of variables may be hardened
  - A subset of the code may be hardened
- In this way, the most suitable trade-off between detection capabilities and overhead can be attained



# Extension to fault tolerance

- The rules can be extended to obtain fault tolerance by means of forward error correction
- Only faults affecting data have been targeted



# SIHFT techniques

- Instruction-level time redundancy
  - Data-oriented techniques
  - Control-oriented techniques
- Task-level time redundancy



# Control-oriented techniques

- Aim at detecting faults changing the execution flow of a program
- All instruction-level techniques are based on:
  - Dividing the program code in **basic blocks**
  - Building the **program graph**
  - Checking at run-time the correctness of each transition performed during the program execution



# Basic blocks (BBs)

- A **basic block** (also called *branch free interval*) is a maximal sequence set of consecutive program instructions that, in absence of faults, are always executed together from the first to the last one
- There is no branching instruction in a BB except possibly for the last one
- No instructions within the BB can be the destination of a branch, jump or call instruction, except for the first one, possibly



# Program Graph (PG)

- It is a graph where:
  - Vertices are basic blocks
  - Edges are transitions in the execution flow, i.e., branches, call and return instructions, etc.
- Let's call  $V$  the set of nodes, and  $B$  the set of edges in PG



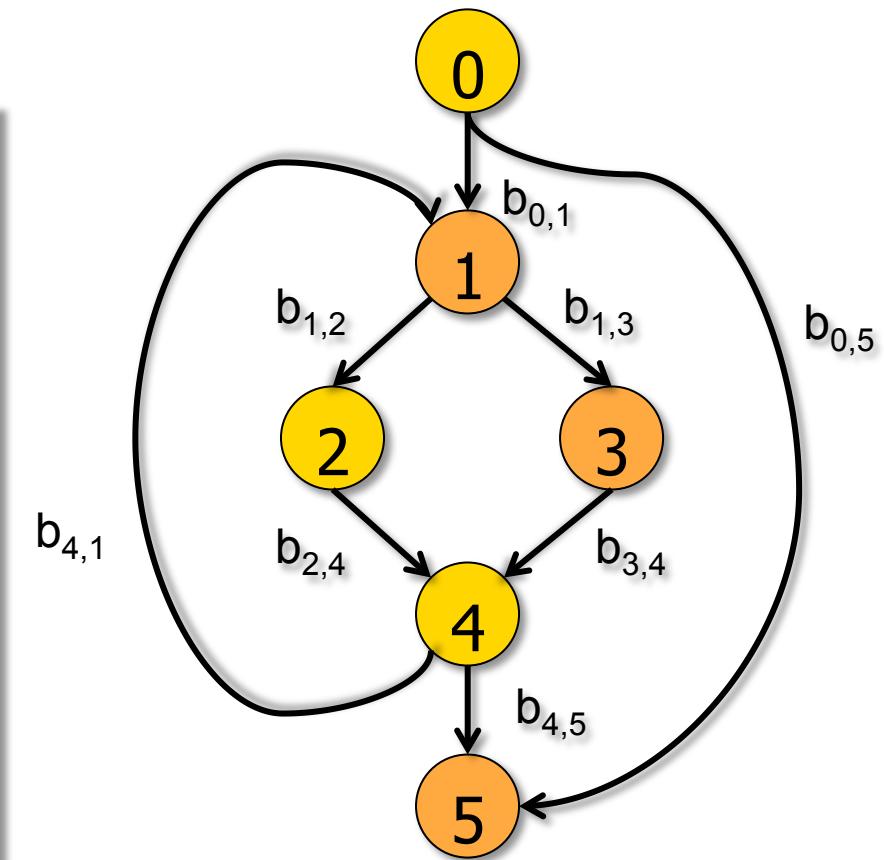
# Example

```
0:    i = 0;
0:    while( i < n ) {
1:        if( a[i] < b[i] )
2:            x[i] = a[i];
3:        else
3:            x[i] = b[i];
4:        i++; }
5:
```



# Example

```
0:    i = 0;  
0:    while( i < n ) {  
1:        if( a[i] < b[i] )  
2:            x[i] = a[i];  
3:        else  
3:            x[i] = b[i];  
4:        i++; }  
5:
```



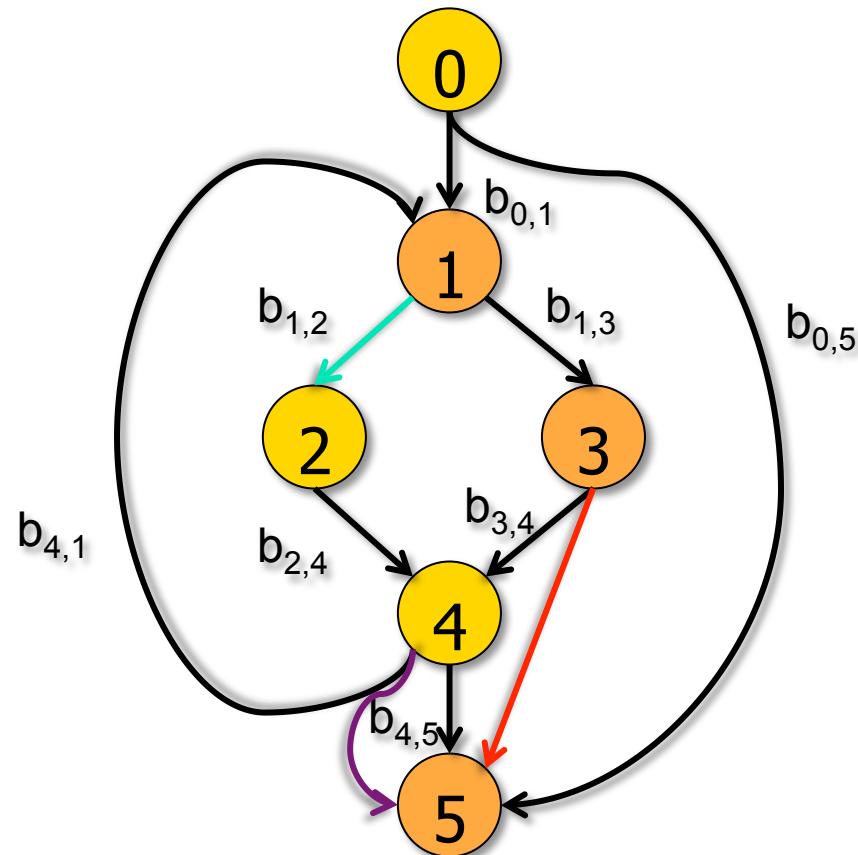
# Edge classification

- An edge connecting two vertices in  $V$  can be:
  - Legal, if it belongs to  $B$
  - Illegal, if it does not belong to  $B$
  - Wrong, if it is legal, but not consistent with the input data
- Illegal and wrong edges correspond to Control Flow Errors (CFEs)



# Examples

- Legal
- Illegal
- Wrong



# Control Flow Errors

- They include
  - Wrong and illegal edges in the Program Graph
  - Edges that can not be modeled in the Program Graph  
(e.g., CFEs causing a wrong branch within a BB)



# Control Flow Error causes

- They may include
  - Faults in the Instruction Register
  - Faults in the Program Counter
  - Faults in the stack
  - Faults in the processor decode unit
  - ...



# Control Flow Checking

- Ideally, it can be implemented by:
  - Building the Program Graph
  - Monitoring the program execution
  - Checking whether the followed control flow is compatible with the Program Graph



# Control Flow Checking

- Ideally, it can be implemented by:
  - Building the Program Graph
  - Monitoring the program execution using a block identifier
  - Checking whether the followed control flow is compatible with the Program Graph

When the SIHFT approach is followed, this is done by embedding in the original code some additional instruction



# Selected approaches

- ECCA (Abraham et al.)
- CFCSS (McCluskey et al.)
- YACCA (Goloubeva et al.)

# Selected approaches

- ECCA (Abraham et al.)
- CRA (McCluskey et al.)

- A *test assertion* is executed at the beginning of the block and checks if the previous basic block is permissible
- A *set assignment* is executed at the end of the block and updates a unique block identifier



# Selected approaches

- ECCA (Abraham et al.)
- CFCSS (McCluskey et al.)
- YACC (Yannakakis et al.)

At the beginning of the block:

- A set assignment updates a unique block identifier
- A test assertion checks the correctness of the identifier



# Selected approaches

At the beginning AND at the end of each basic block additional instructions are introduced:

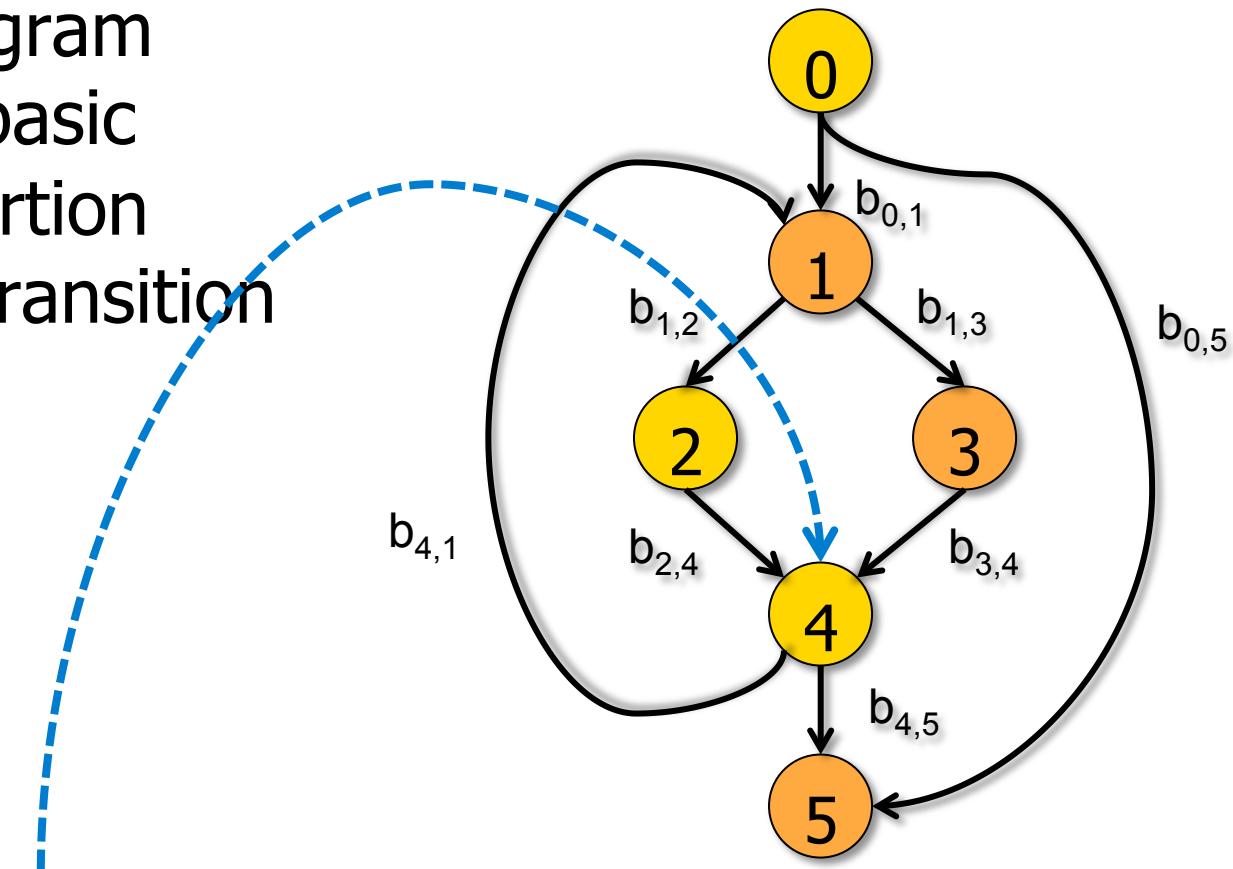
- A test assertion controls the current signature and checks if it is permissible
- A set assignment updates the signature

- YACCA (Goloubeva et al.)



# Test Assertion

- When the program enters into a basic block, an assertion checks if the transition is legal



```
assert( (code==B2) || (code==B3) )
```

# Set Assignment

- Ideally, it could be an assertion such as:

```
code = Bi
```

- Any fault resulting in a jump to this instruction would not be detected



# Set Assignment

- Ideally, it could be an assertion such as:

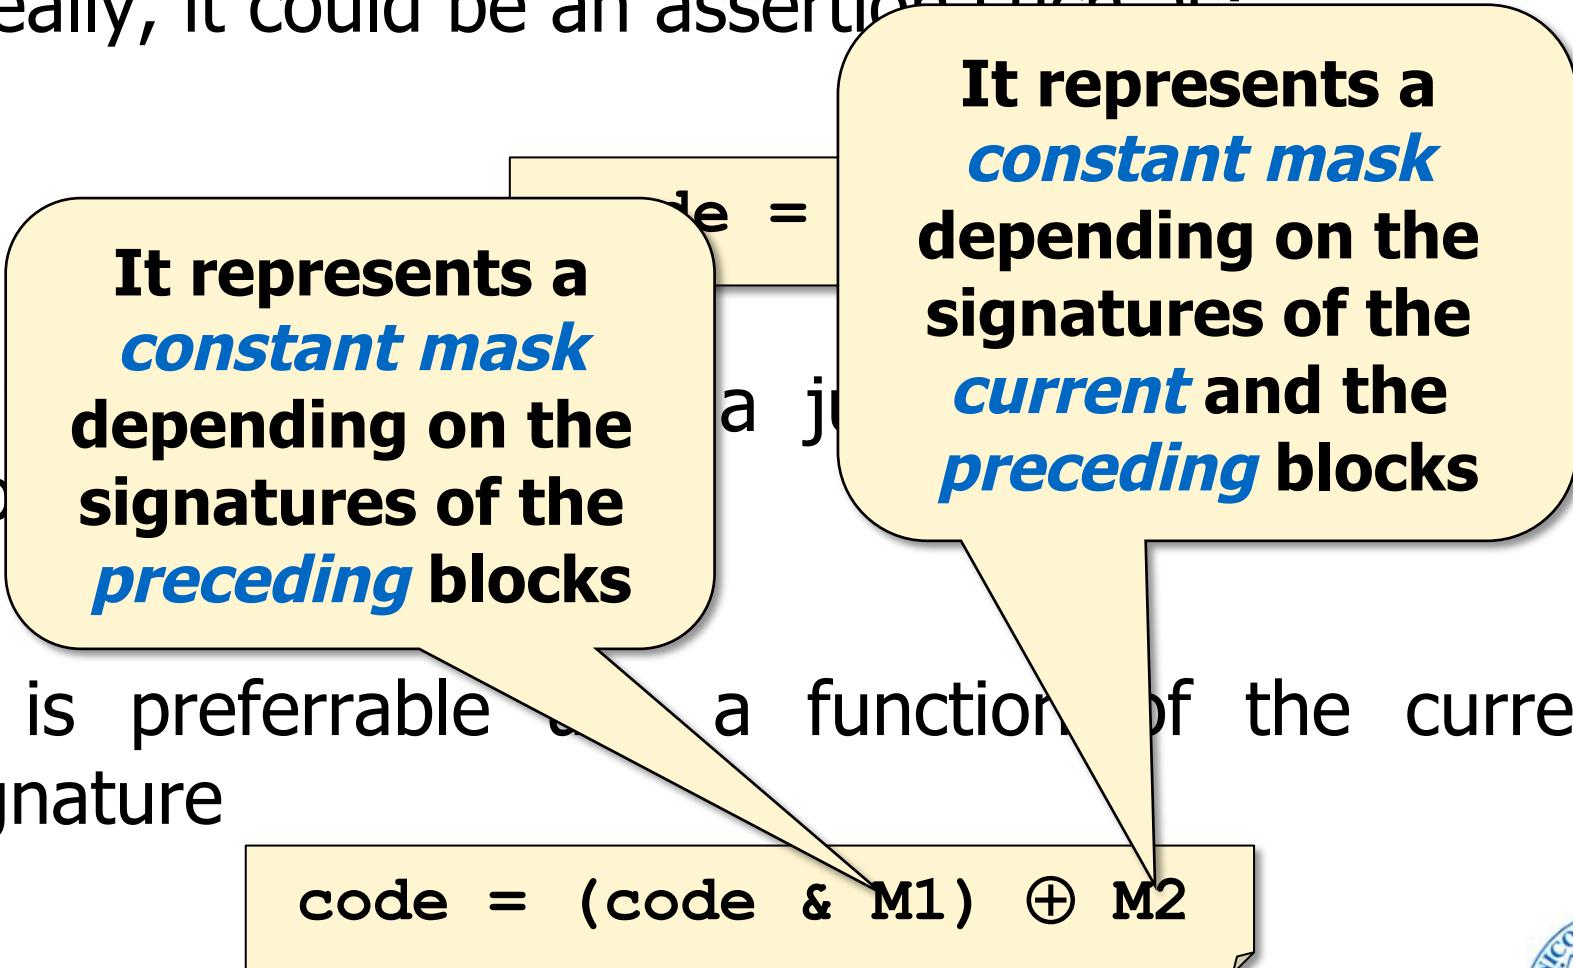
```
code = Bi
```

- Any fault resulting in a jump to this instruction would not be detected
- It is preferable use a function of the current signature

```
code = (code & M1) ⊕ M2
```



# Set Assignment

- Ideally, it could be an assertion such as:  


code = (code & M1) ⊕ M2

It represents a **constant mask** depending on the signatures of the **preceding** blocks

It represents a **constant mask** depending on the signatures of the **current** and the **preceding** blocks
- An example would be:
- It is preferable to use a function of the current signature

# Set Assignment

- Ideally, it could be an assertion such as:  
**Both M1 and M2 can be computed for each assertion at compile time**
- Any fault would result in this instruction
- It is preferable to use a function of the current signature

```
code = (code & M1) ⊕ M2
```

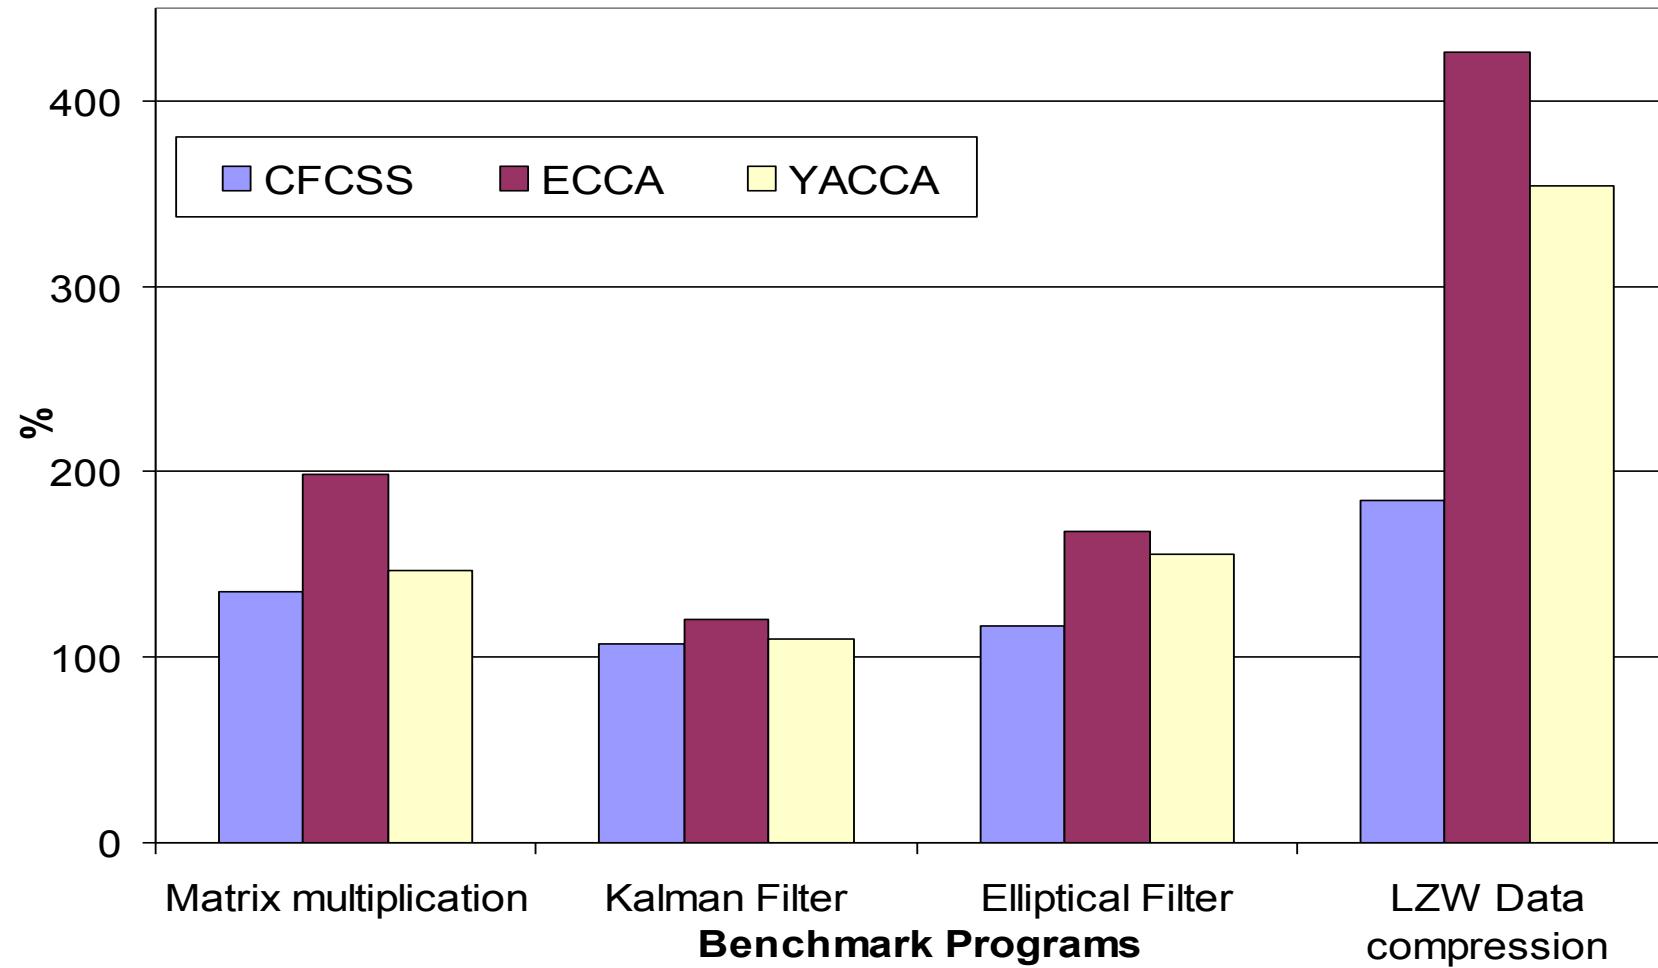


# Experimental environment

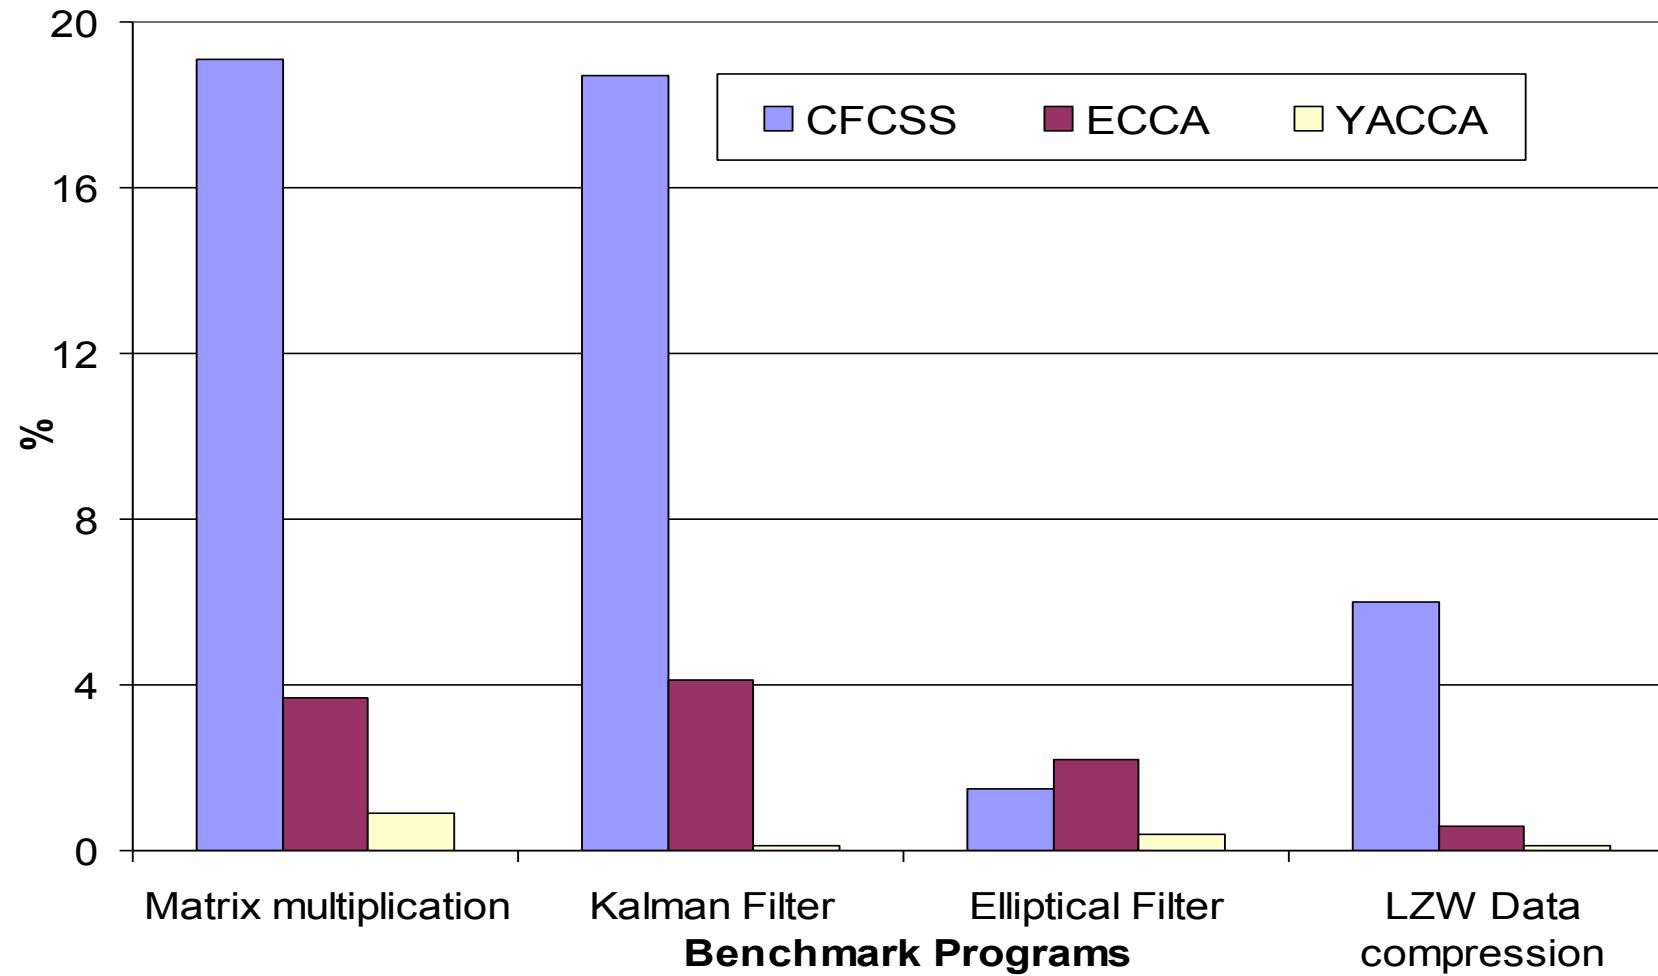
- Simple benchmark programs (in C language)
- Target system based on a SPARC microprocessor
- Code transformations implemented by an automatic tool
- Simulation-based Fault Injection environment



# Performance slow-down



# Wrong answers



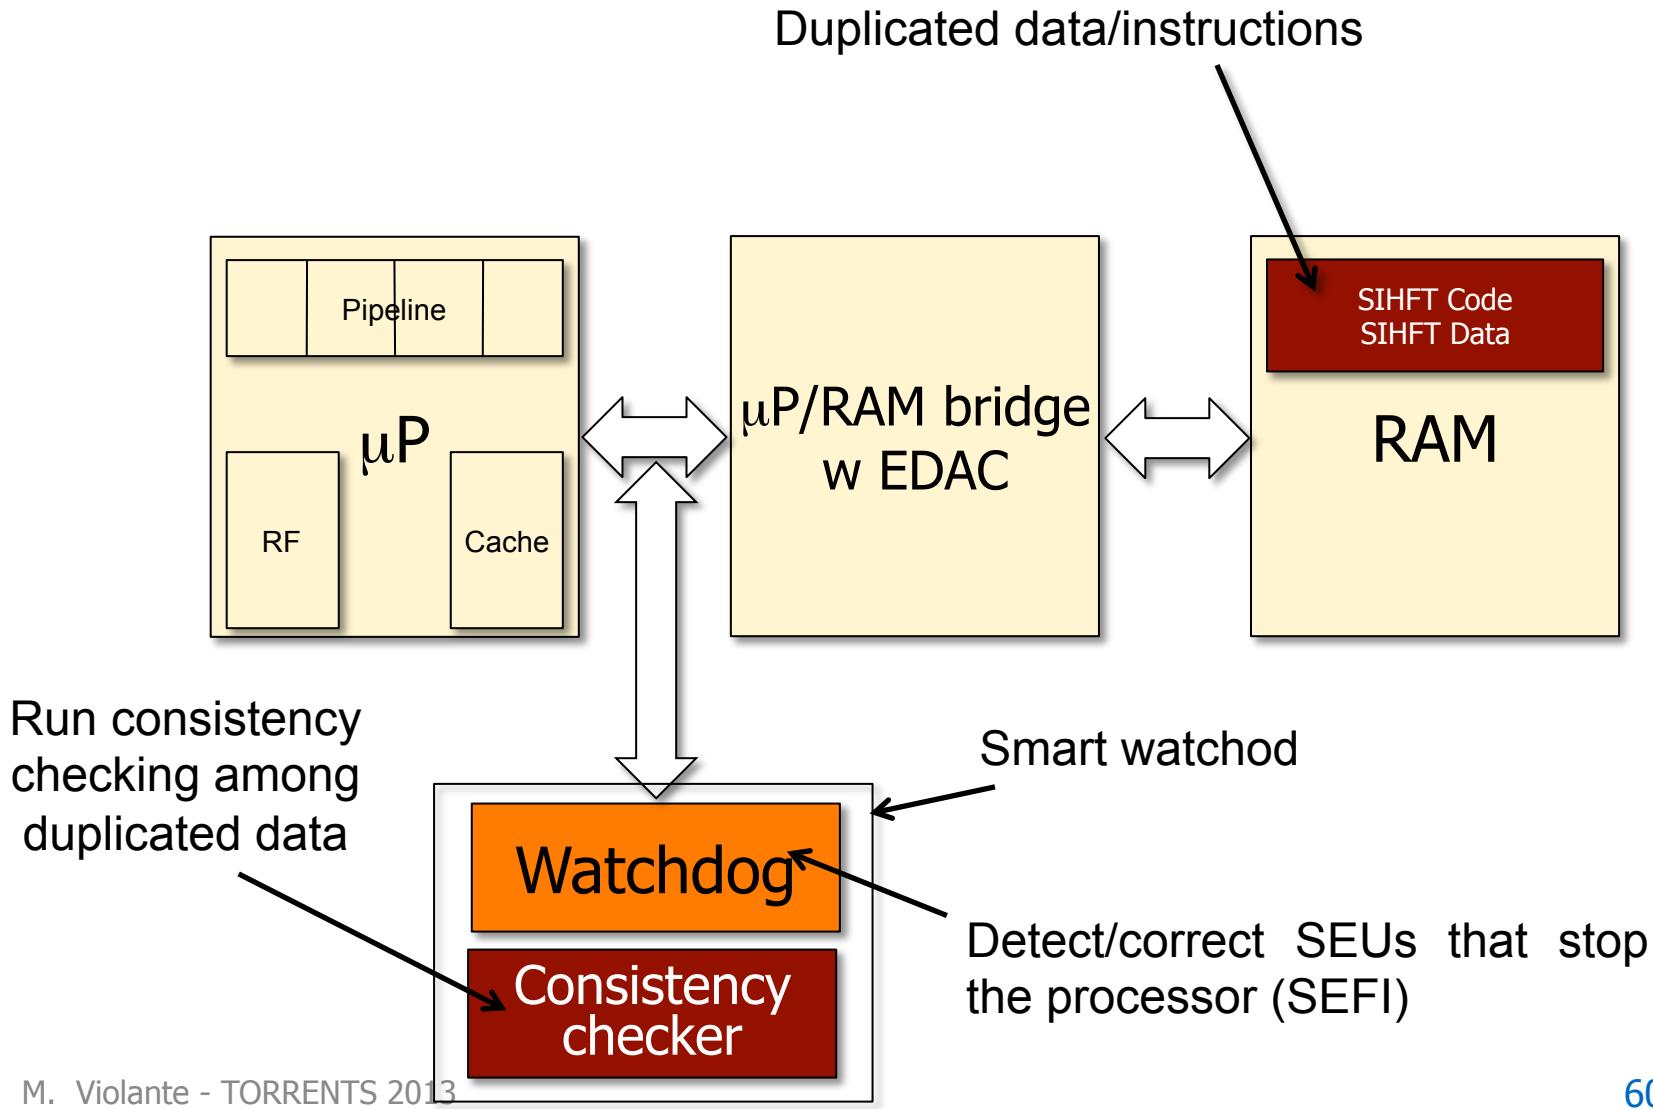
# Coping with performance issues

- Performance overhead due to three factors:
  - Need for disabling compiler optimization → hard to avoid
    - SIHFT methods assumes a certain ordering of instructions
  - Duplication of operations → impossible to avoid
    - SIHFT exploits time redundancy
  - Execution of consistency checks → avoidable?

```
int a1, b1, a2, b2;  
...  
b1 = a1+5;          b2 = a2+5;  
assert( a1==a2 );
```



# Improved System architecture



# Hardware consistency checker

- Exploits specific ordering of duplicated data obtained when compiling the SIHFT software
  - $\text{ADX}(a1) = \text{ADX}(a0)+\text{OFFSET}$
- “Snoops” read/write cycles and stores (ADX, value) pairs in a context addressable memory (CAM)
  - ADX is used as key for the CAM
- Based on ADX value, it identifies replicas of the same data and runs consistency checks



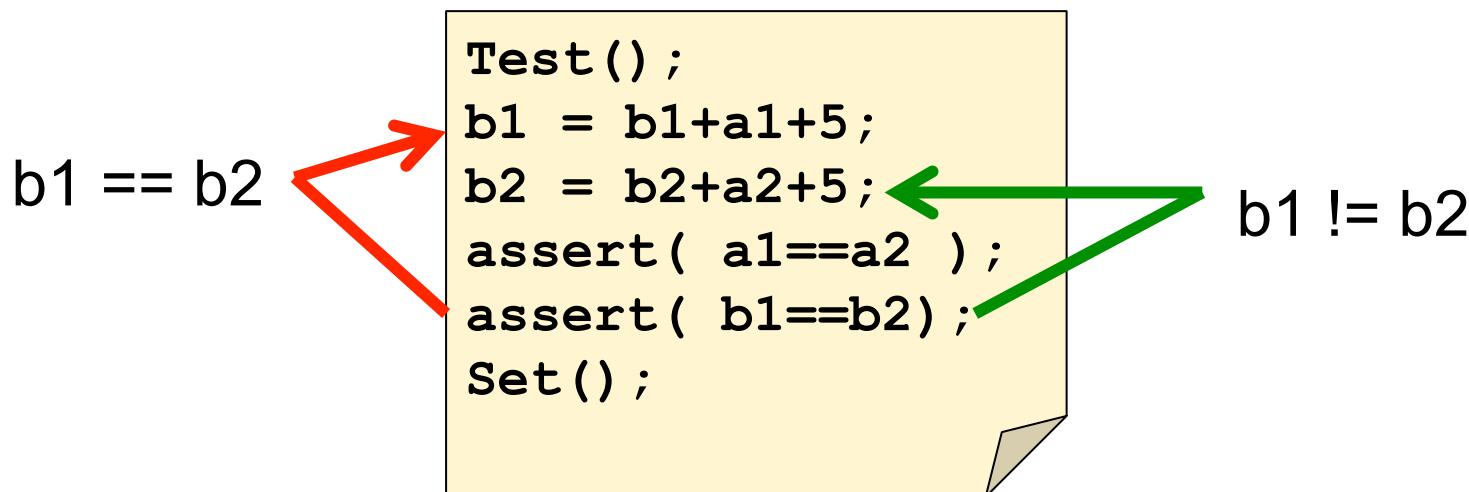
# Hardware consistency checker

- Same fault coverage
- Performance overhead  $\sim 2x$  (compared to  $3x-4x$ )
- Area overhead  $< 10\%$



# Wrap-up

- Instruction-level SIHFT detects:
  - 100% SEE affecting data (no matter where it is stored)
  - >98% SEE affecting control flow
    - SEE provoking jumps within the basic block can escape



- Costs: >2x performance, < 10% area

# SIHFT techniques

- Instruction-level time redundancy
  - Data-oriented techniques
  - Control-oriented techniques
- Task-level time redundancy



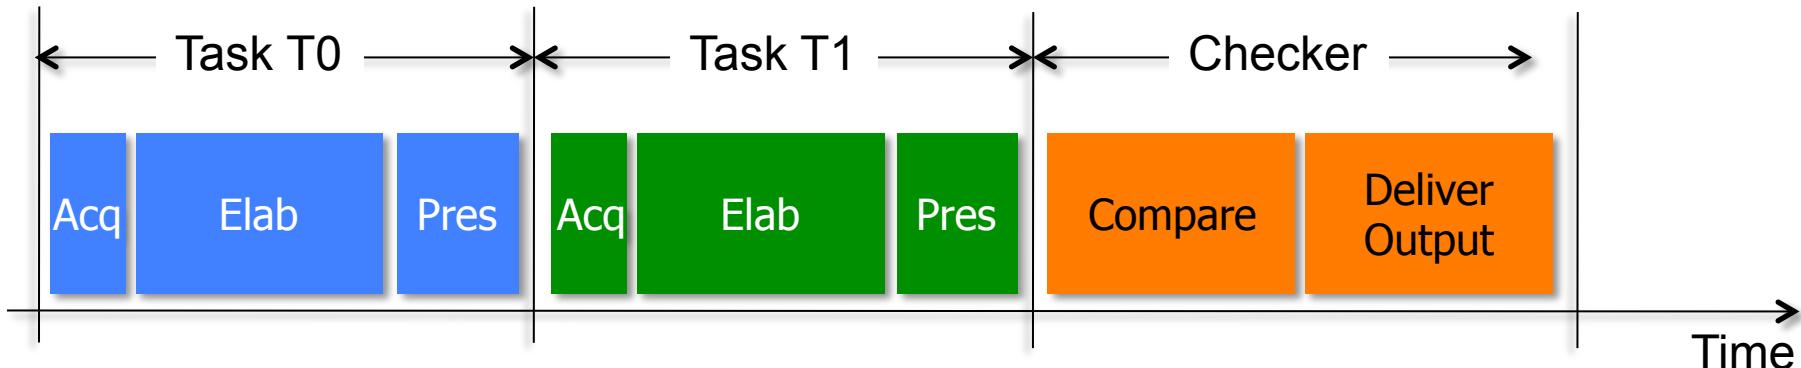
# Task-level redundancy

- Idea:
  - Time redundancy is applied at task level
  - Consistency checks done over task outputs
- Lower performance overhead:
  - Fewer consistency checks
  - No need for disabling compiler optimizations
  - Applicable also when source code not available (e.g., libraries)
- Require substantial “manual” design
  - No automation



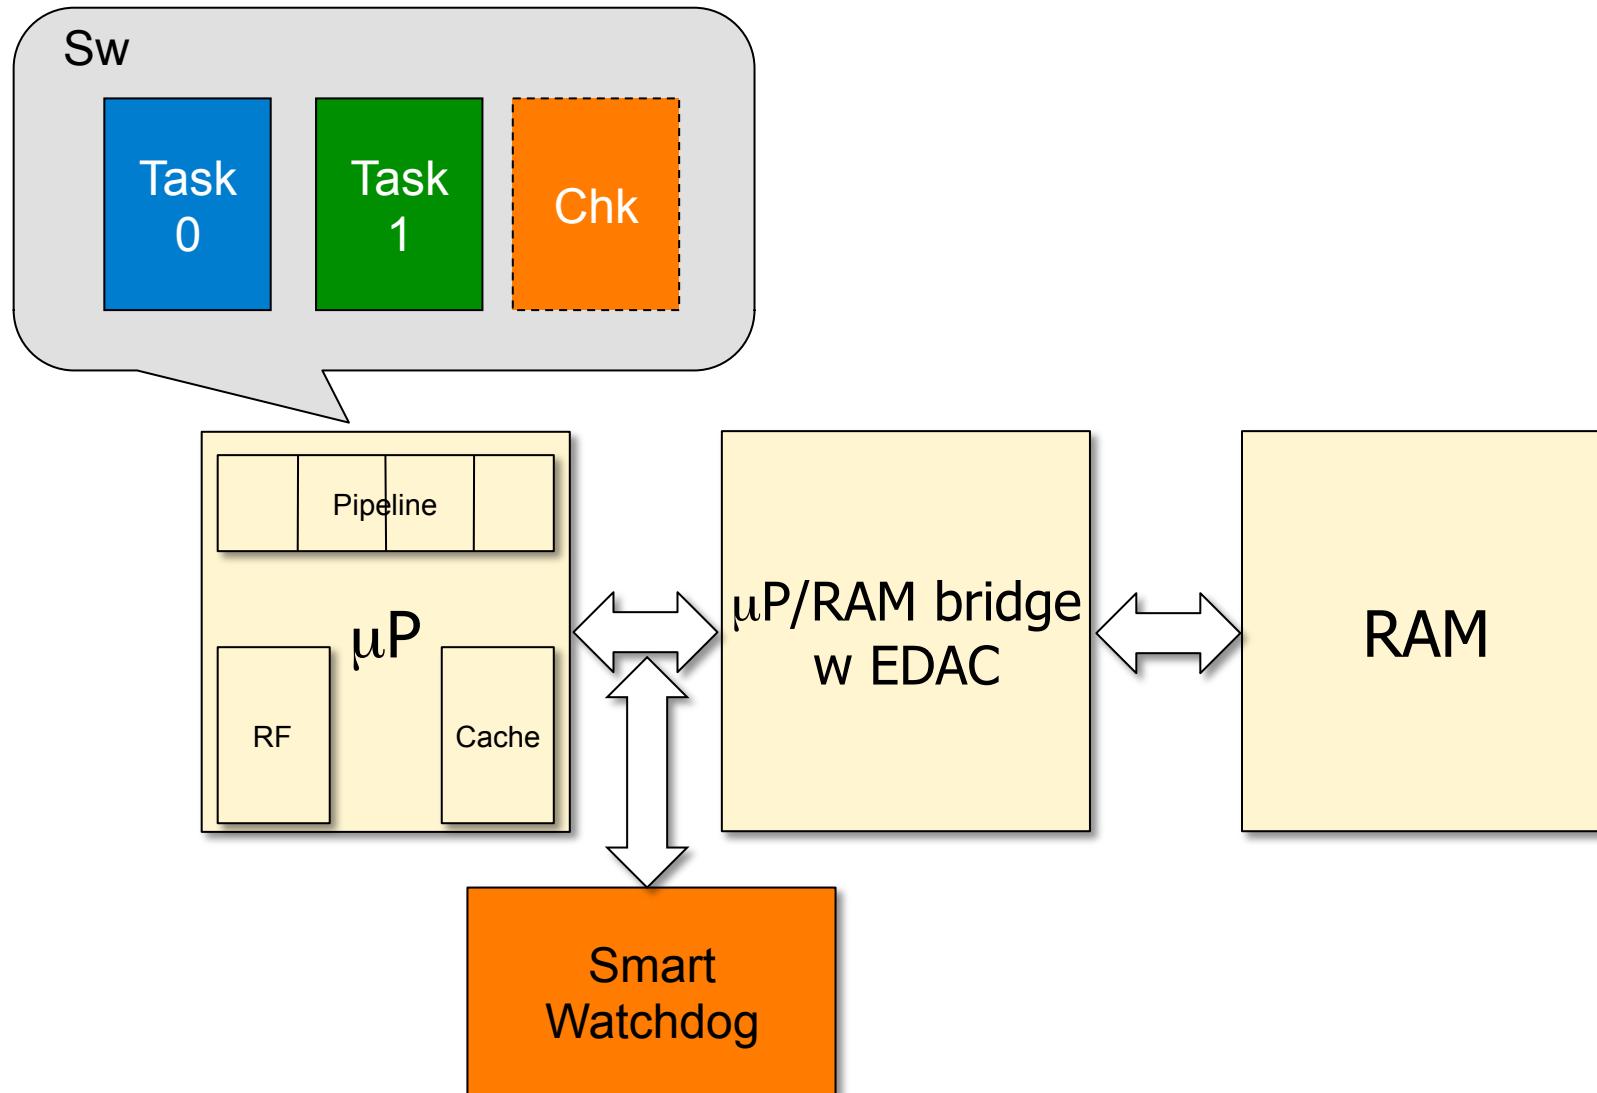
# Task-level redundancy in a nutshell

- Run two instances ( $T_0, T_1$ ) of the same task

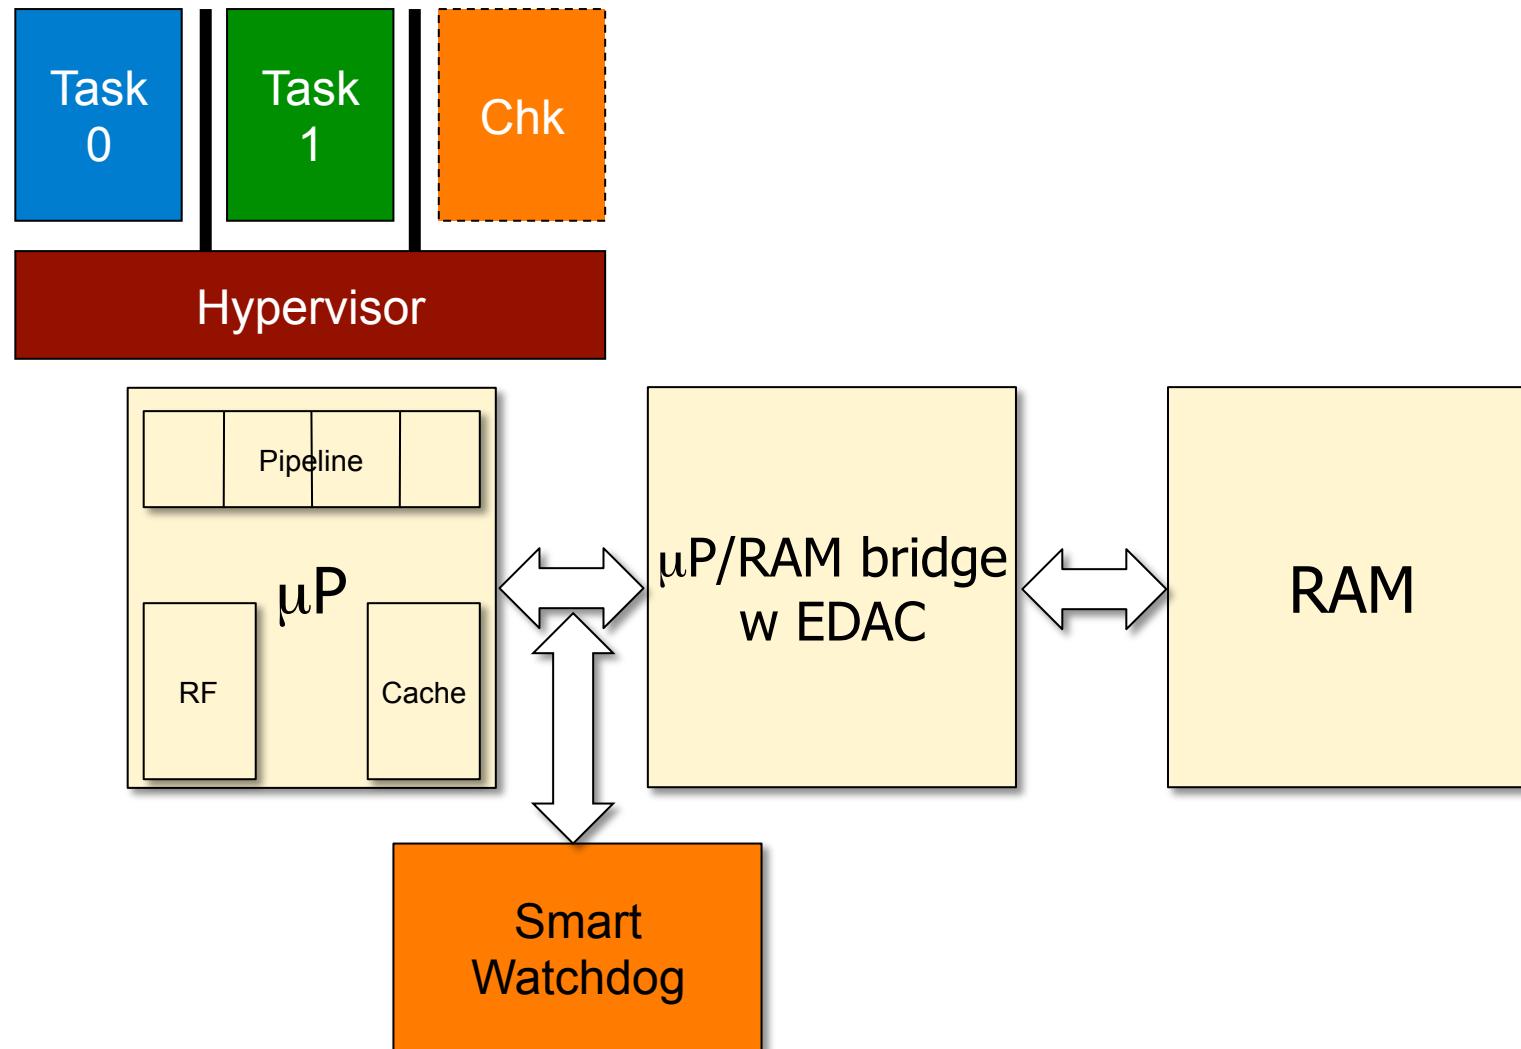


- The executions are segregated in their own address space
  - Task T0 should never corrupt Task T1
- “Independent” consistency check detects errors
  - Run by custom sw or hw

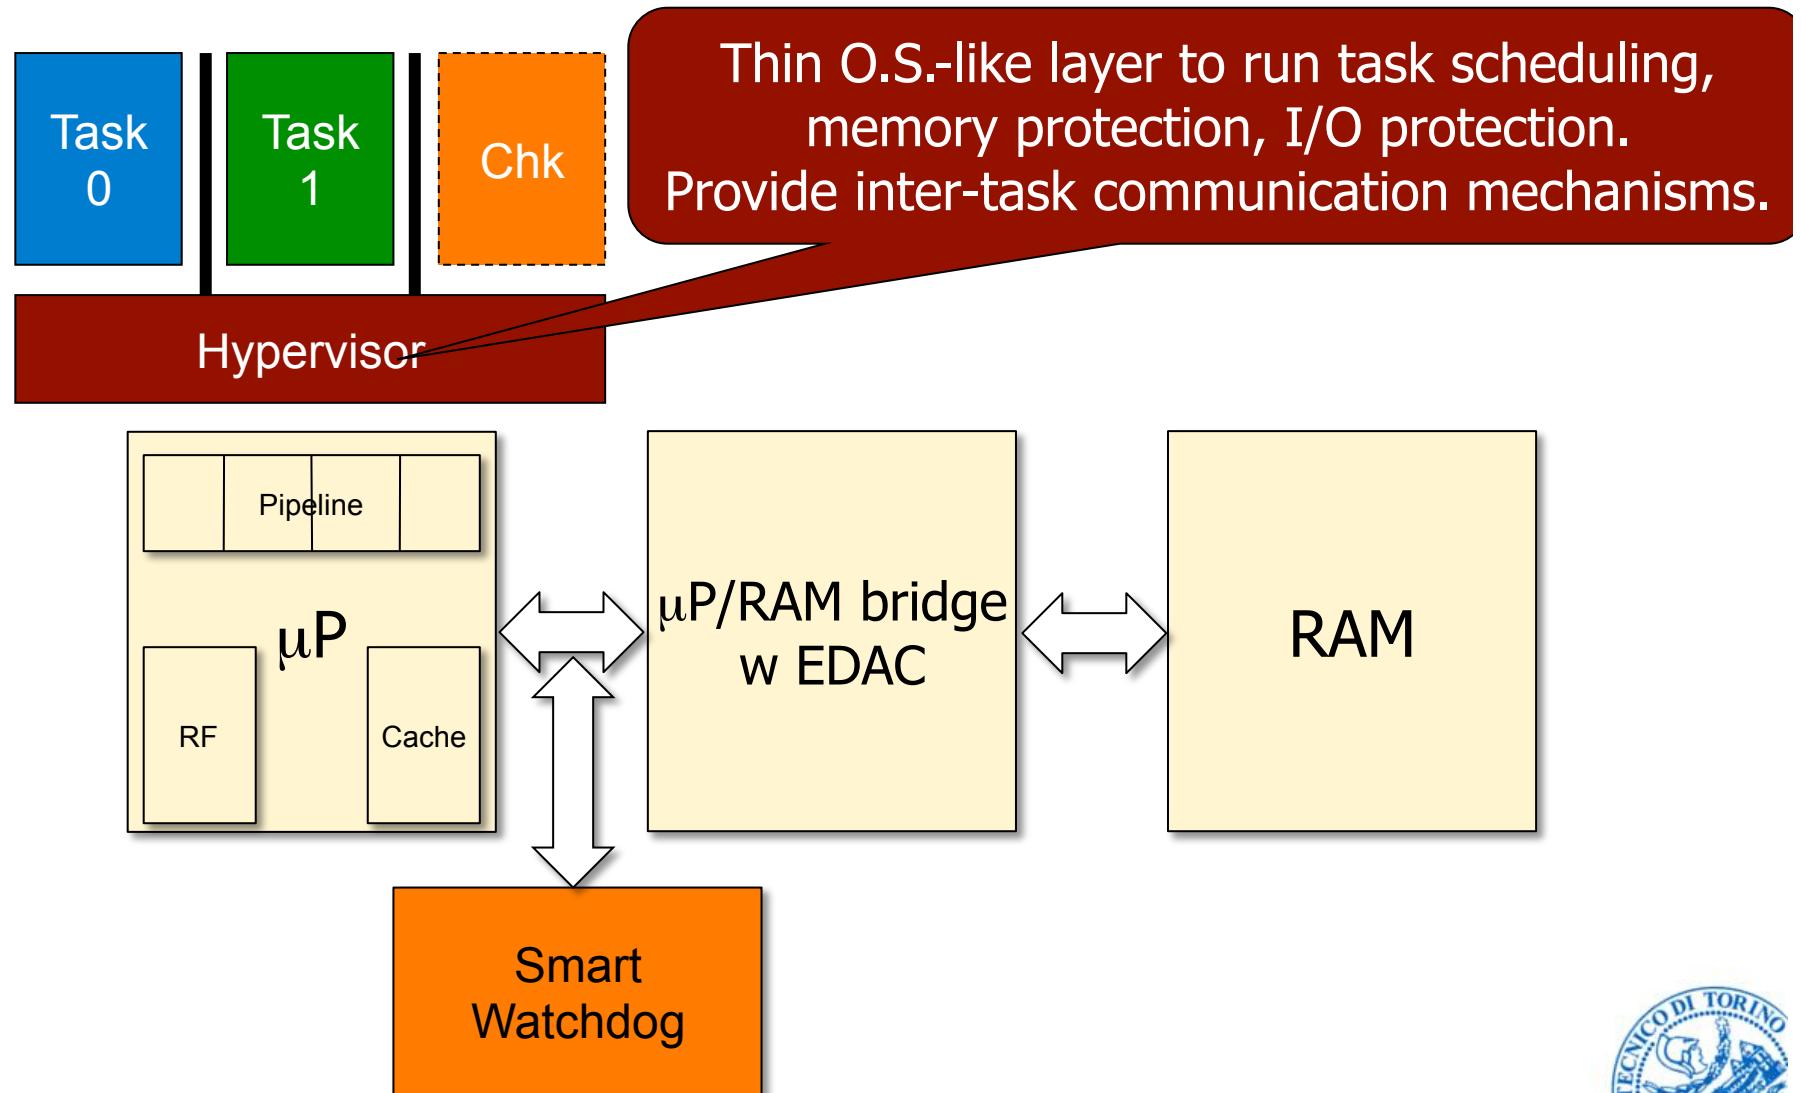
# General task-level redundancy



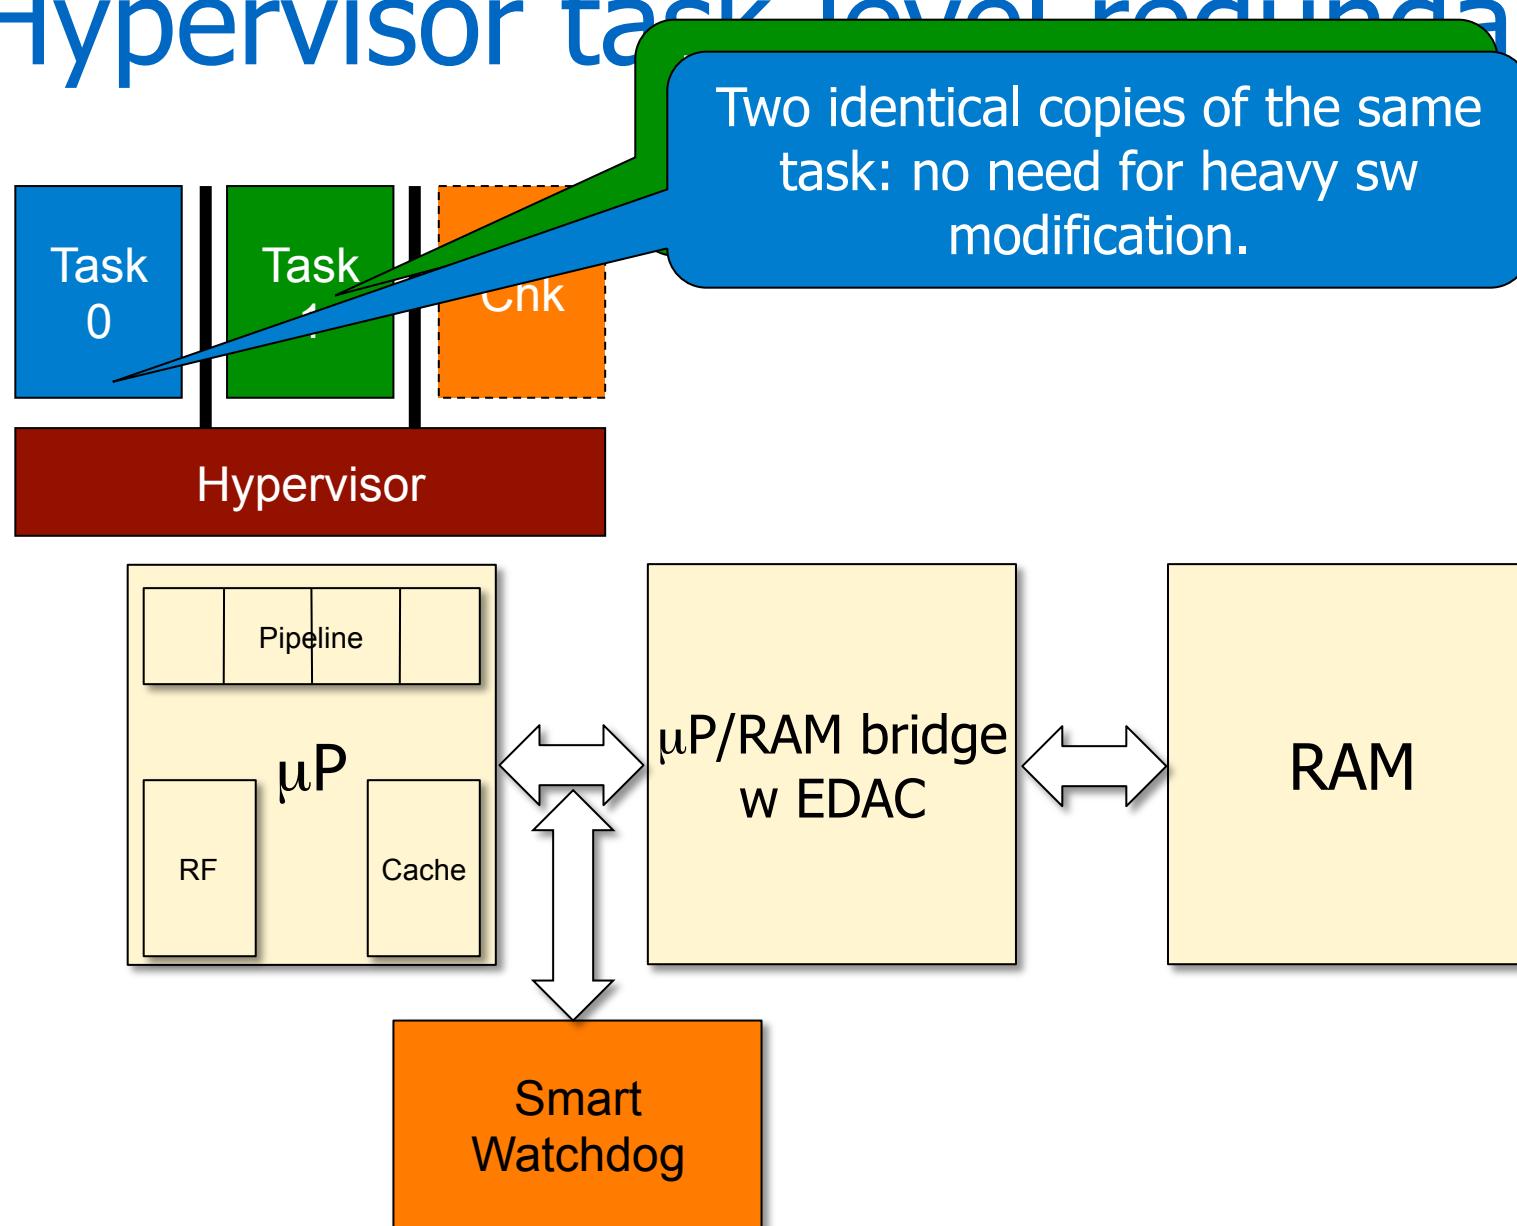
# Hypervisor task-level redundancy



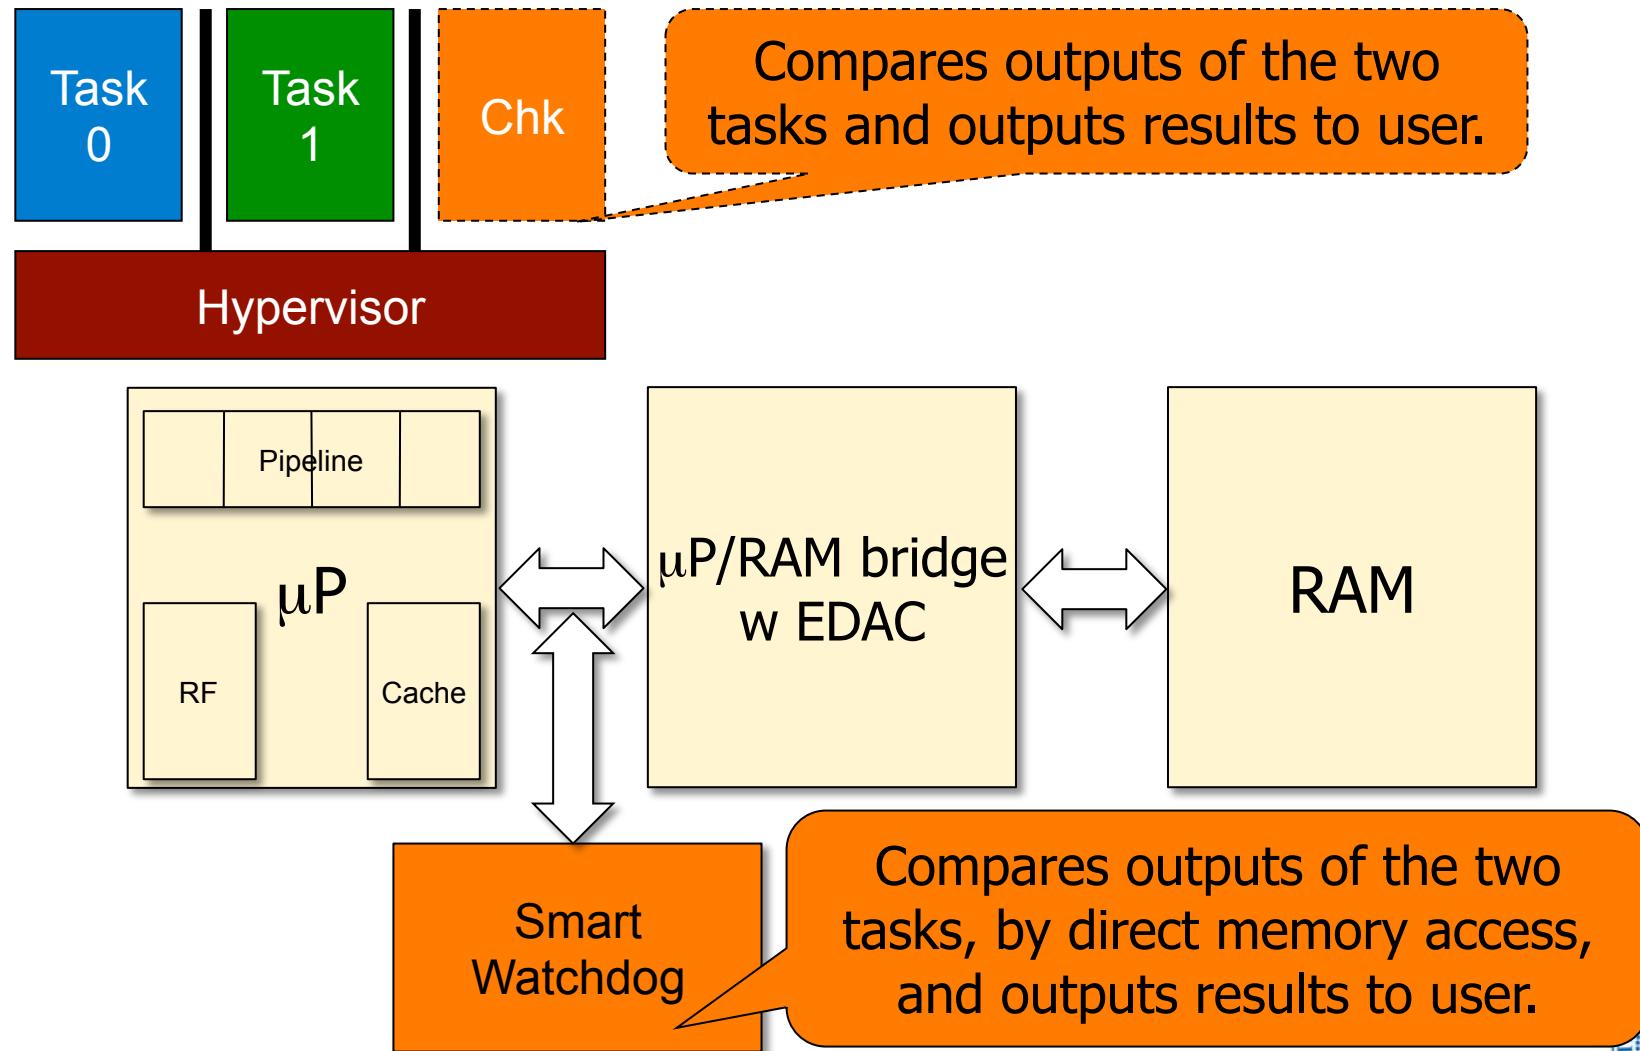
# Hypervisor task-level redundancy



# Hypervisor task level redundancy



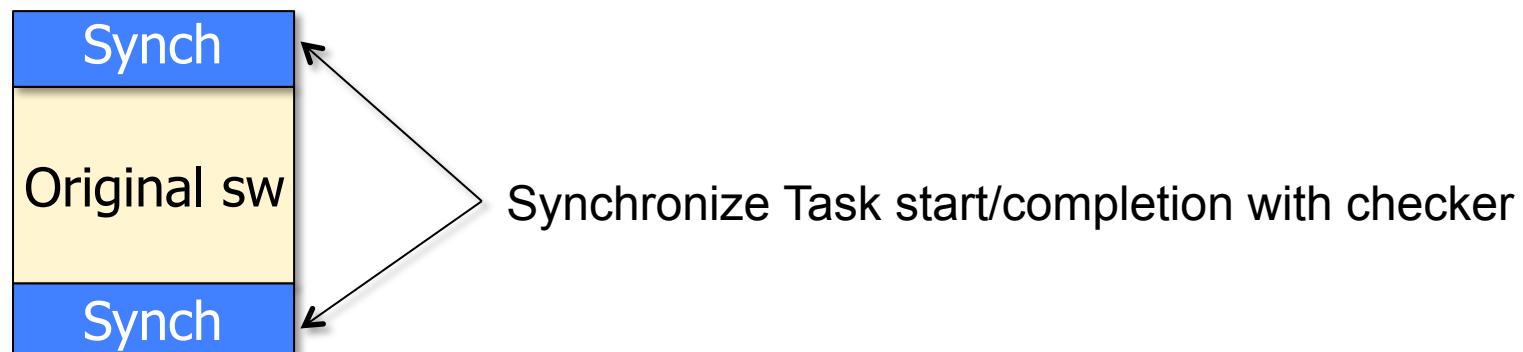
# Hypervisor task-level redundancy



# Advantages/Limitations

## ■ Advantages:

- Limited development effort (most of the complexity dealt by hypervisor, automated)



- Easily portable/scalable (single- vs multi-core)

## ■ Limitations:

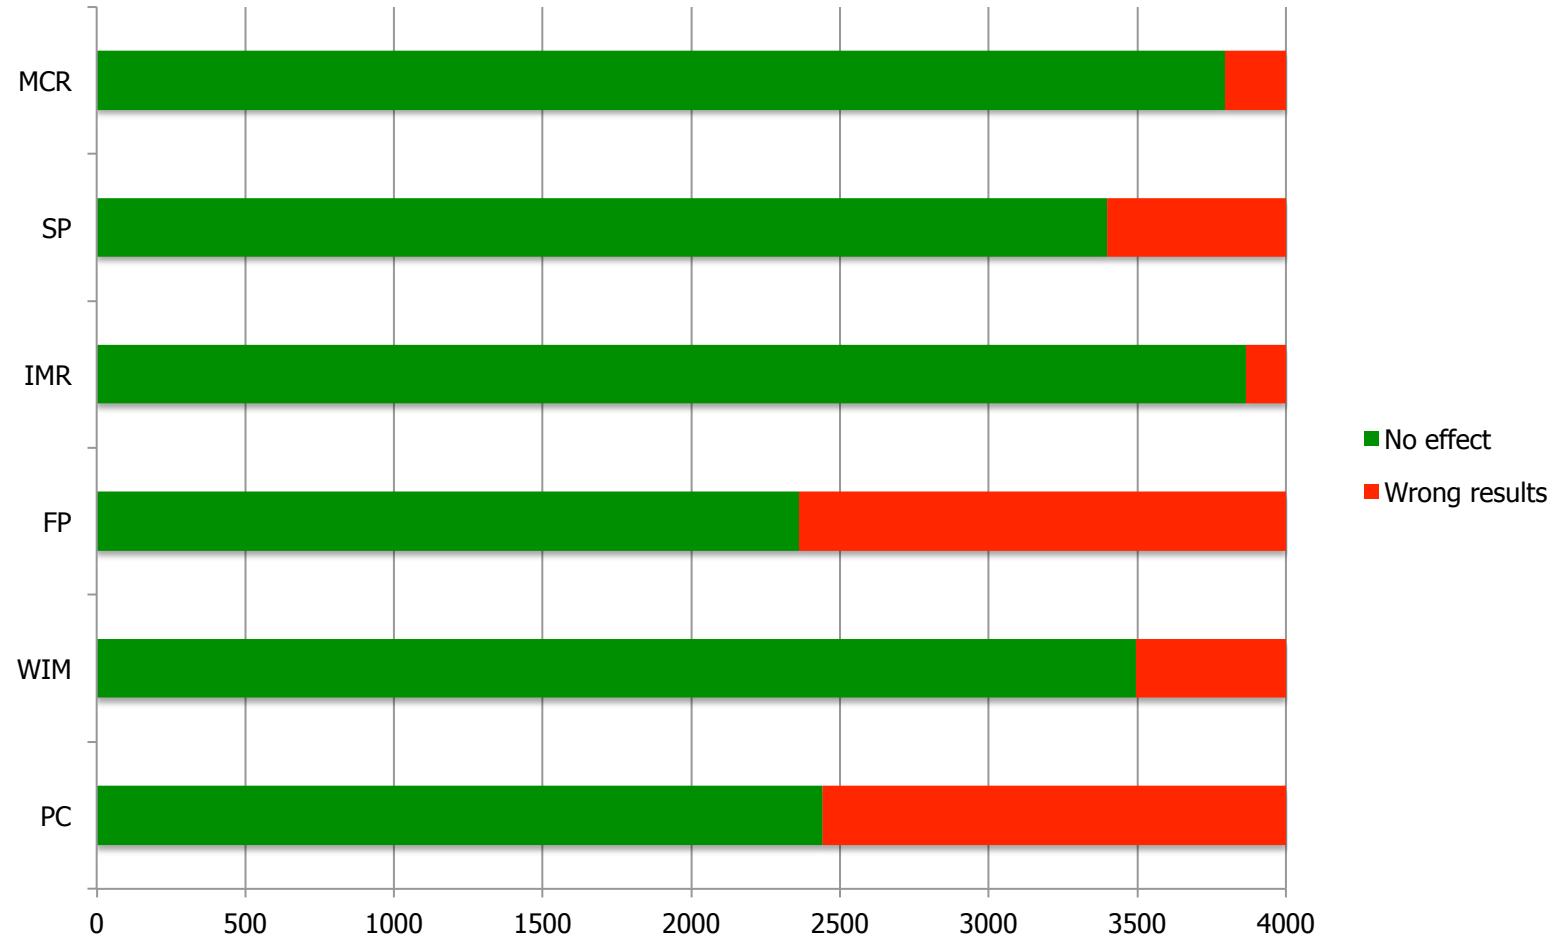
- Hypervisor is not replicated → single point of failure
- Need for hypervisor for the adopted processor

# An example

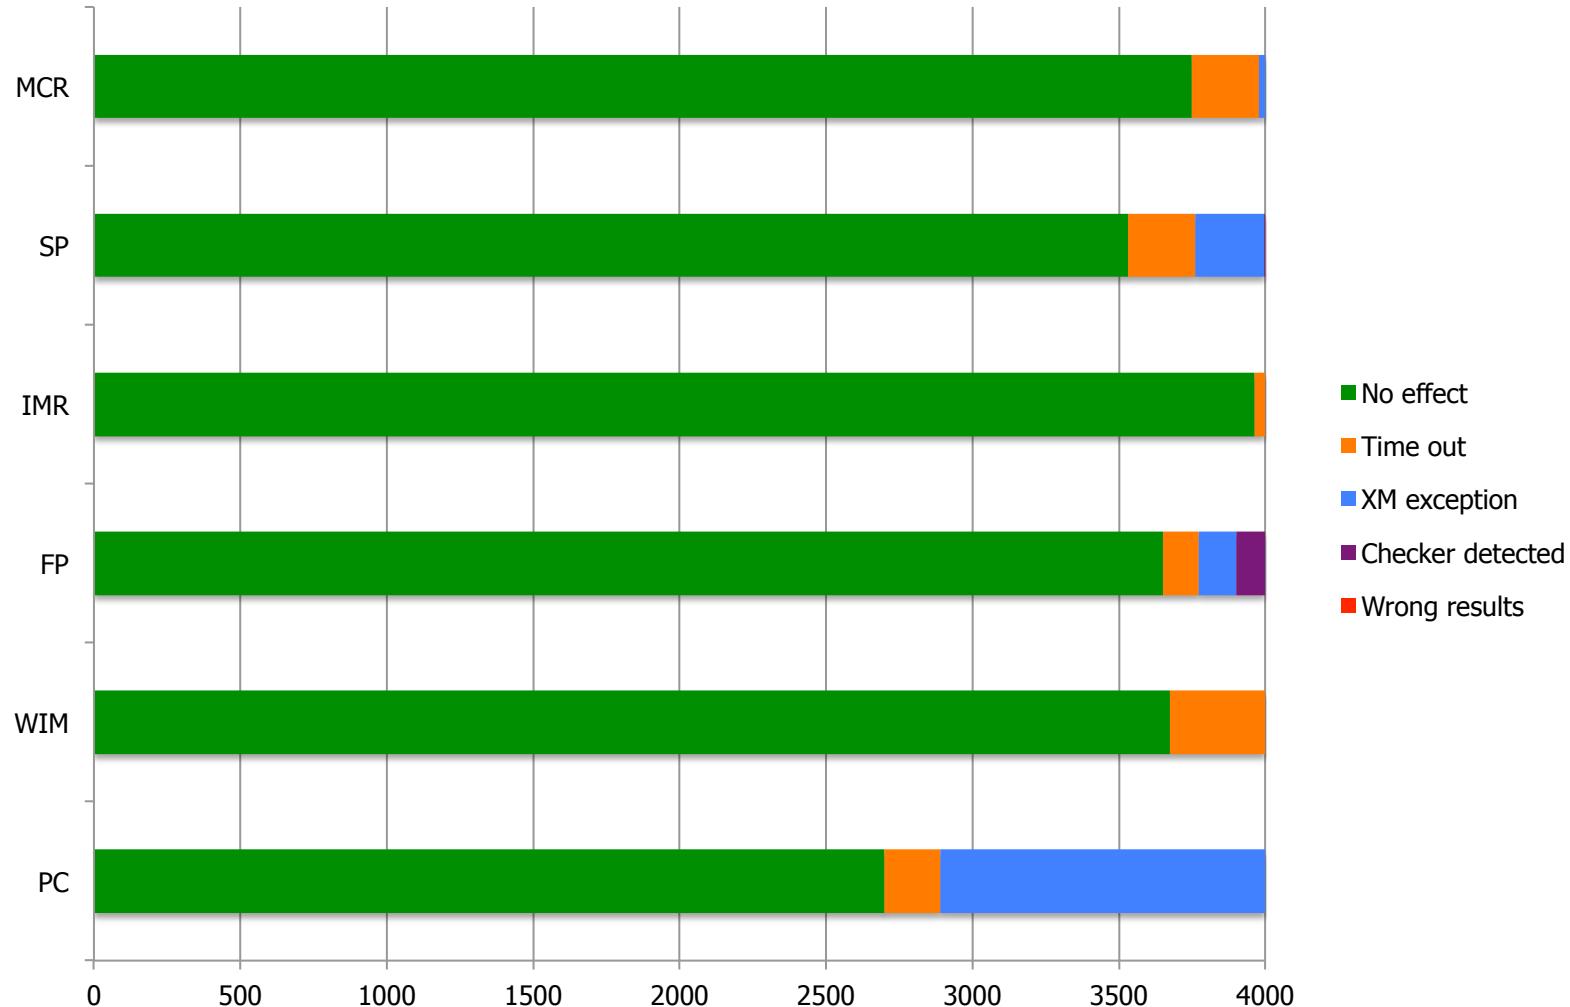
- LEON3 processor from Aeroflex Gaisler
- GR-XC3S-1500 board from Pender
- Xtratum from Universidad Politécnica de Valencia (Spain)
- Smart watchdog coded in VHDL
  - DMA controller (AMBA bus master)
  - Watchdog timer for SEFI detection
  - 3% are overhead w.r.t. the LEON3 core
- Hypervisor vulnerability < 3% execution time



# Original system



# Hardened system



# Conclusions

- SIHFT is viable solution for SEE mitigation in COTS processor (especially for payload processing)
- Instruction-level techniques seems best fit for DSP-based applications where ad-hoc compiler can be used
  - Plenty of execution slots can run duplicated instruction in parallel
- Task-level techniques seems best fit for general-purpose processors
  - Hypervisor can greatly help

